

Cryptography V: Key Exchange, Public-Key Cryptography, and Digital Signature

CSE 565: Fall 2024

Computer Security

Xiangyu Guo (xiangyug@buffalo.edu)

Announcement

- Please sign-up at course Piazza.
- Reminder of Quiz 0 (**Due 09/19**).
 - You must obtain full score of the Quiz.
 - Updated so that you can see exactly where you got wrong.
- Assignment 1 & Project 1 has been released (**Due 09/25, 23:59**).

Review of Last Lecture

- MAC: protects integrity (but not confidentiality)
- Hash function: collision resistance
- Authenticated Encryption
 - Encrypt-then-MAC (recommend)
 - MAC-then-Encrypt
- Attacks
 - Padding Oracle
 - Non-atomic decrypt
 - Do not implement A.E. by yourself! Use a standard if possible.

Today's Topic

- Key Exchange: how we establish the shared secret key
 - Diffie-Hellman
- Public Key Cryptography
 - Trapdoor function
 - RSA
 - ~~ElGamal: PKC from Diffie-Hellman~~
 - Digital Signature

Key Exchange

Where are we now?

- Symmetric Key Crypto
 - ▶ Enables confidentiality
 - ▶ Achieved through secret key encryption
 - ▶ Enables authentication and integrity
 - ▶ achieved through MACs and Authenticated Encryption
- In all of the above the sender and receiver **must share a secret key**
 - ▶ Need a secure channel for key distribution
 - ▶ Not possible for parties with no prior relationship

The Diffie-Hellman Protocol

The prime-number field version

Fix a large prime p (e.g. 600 digits ≈ 2000 bits).


Fix an integer g in $\{2, \dots, p\}$

Choose a random
 $a \in \{1, \dots, p-1\}$



"Alice", $A = g^a \mod p$

Choose a random
 $b \in \{1, \dots, p-1\}$



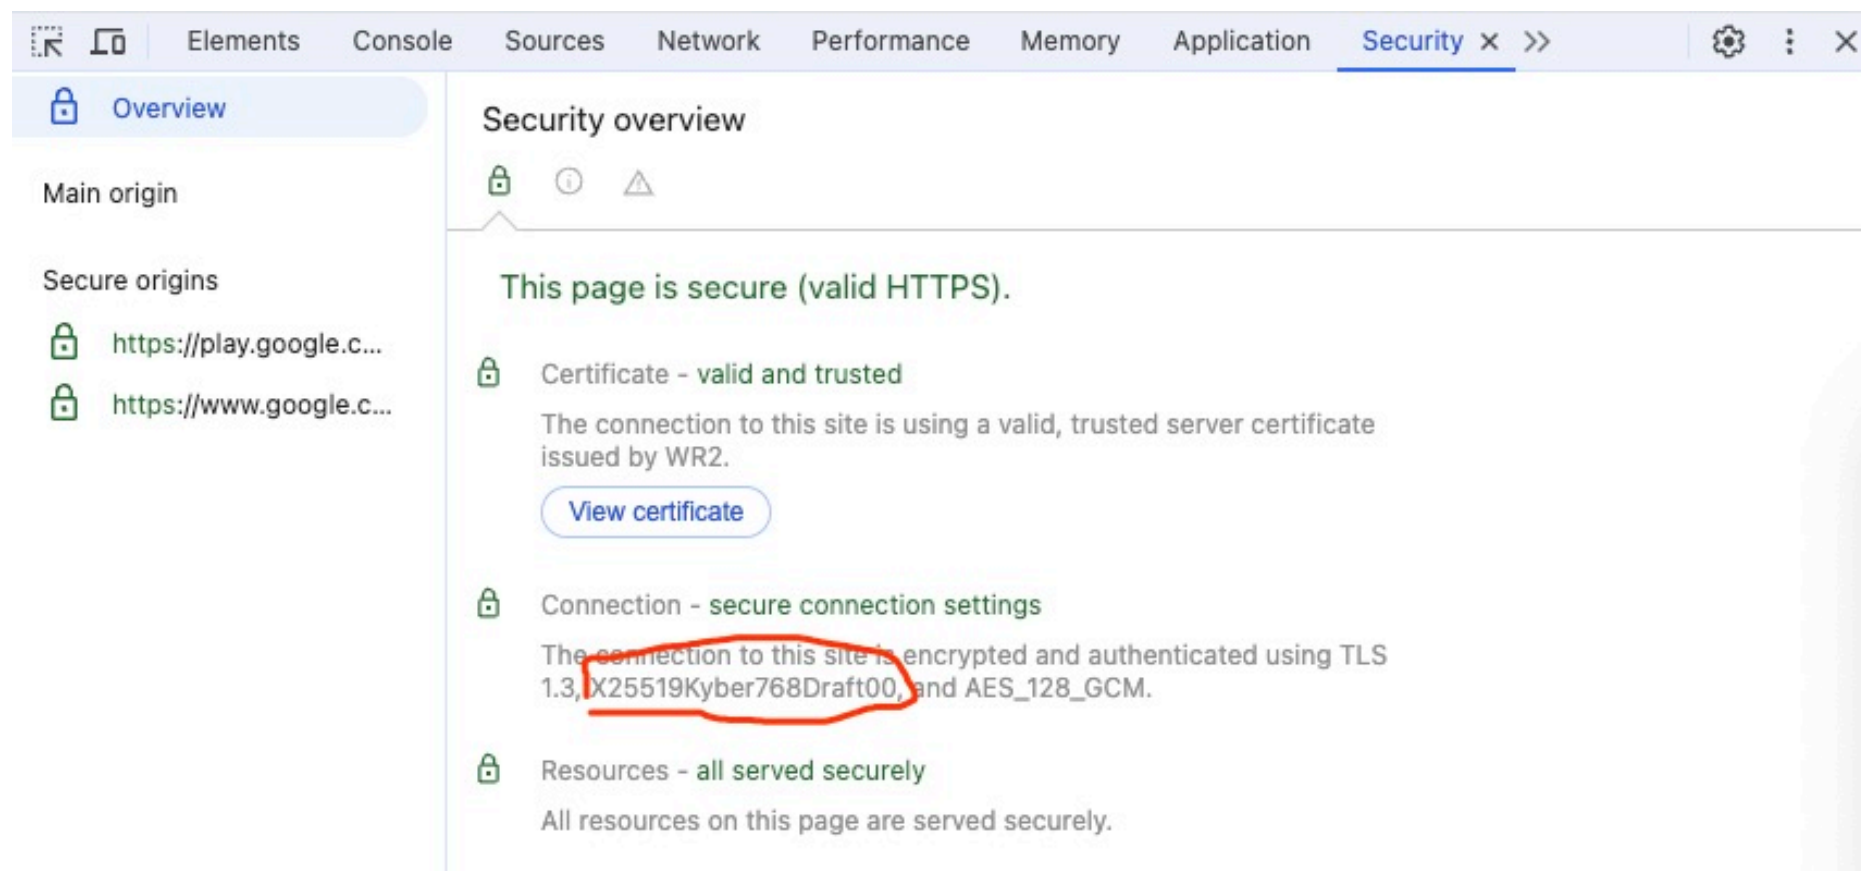
"Bob", $B = g^b \mod p$

$$(g^b)^a = B^a \mod p \quad = \quad k_{AB} = g^{ab} \mod p \quad = \quad (g^a)^b = A^b \mod p$$

shared key

The Diffie-Hellman Protocol

In Chrome's developer tool, you can find the key exchange alg used.



X25519: elliptic curve-based DH key exchange

Kyber768: post-quantum key exchange scheme

Security of Diffie-Hellman

Against Eavesdropping

- Eavesdropper sees: $p, g, A = g^a \bmod p$, and $B = g^b \bmod p$
 - Can compute $g^{ab} \bmod p$?
- More generally, define $\text{DH}_g(g^a, g^b) = g^{ab} \bmod p$
 - How hard is the DH function $\bmod p$?

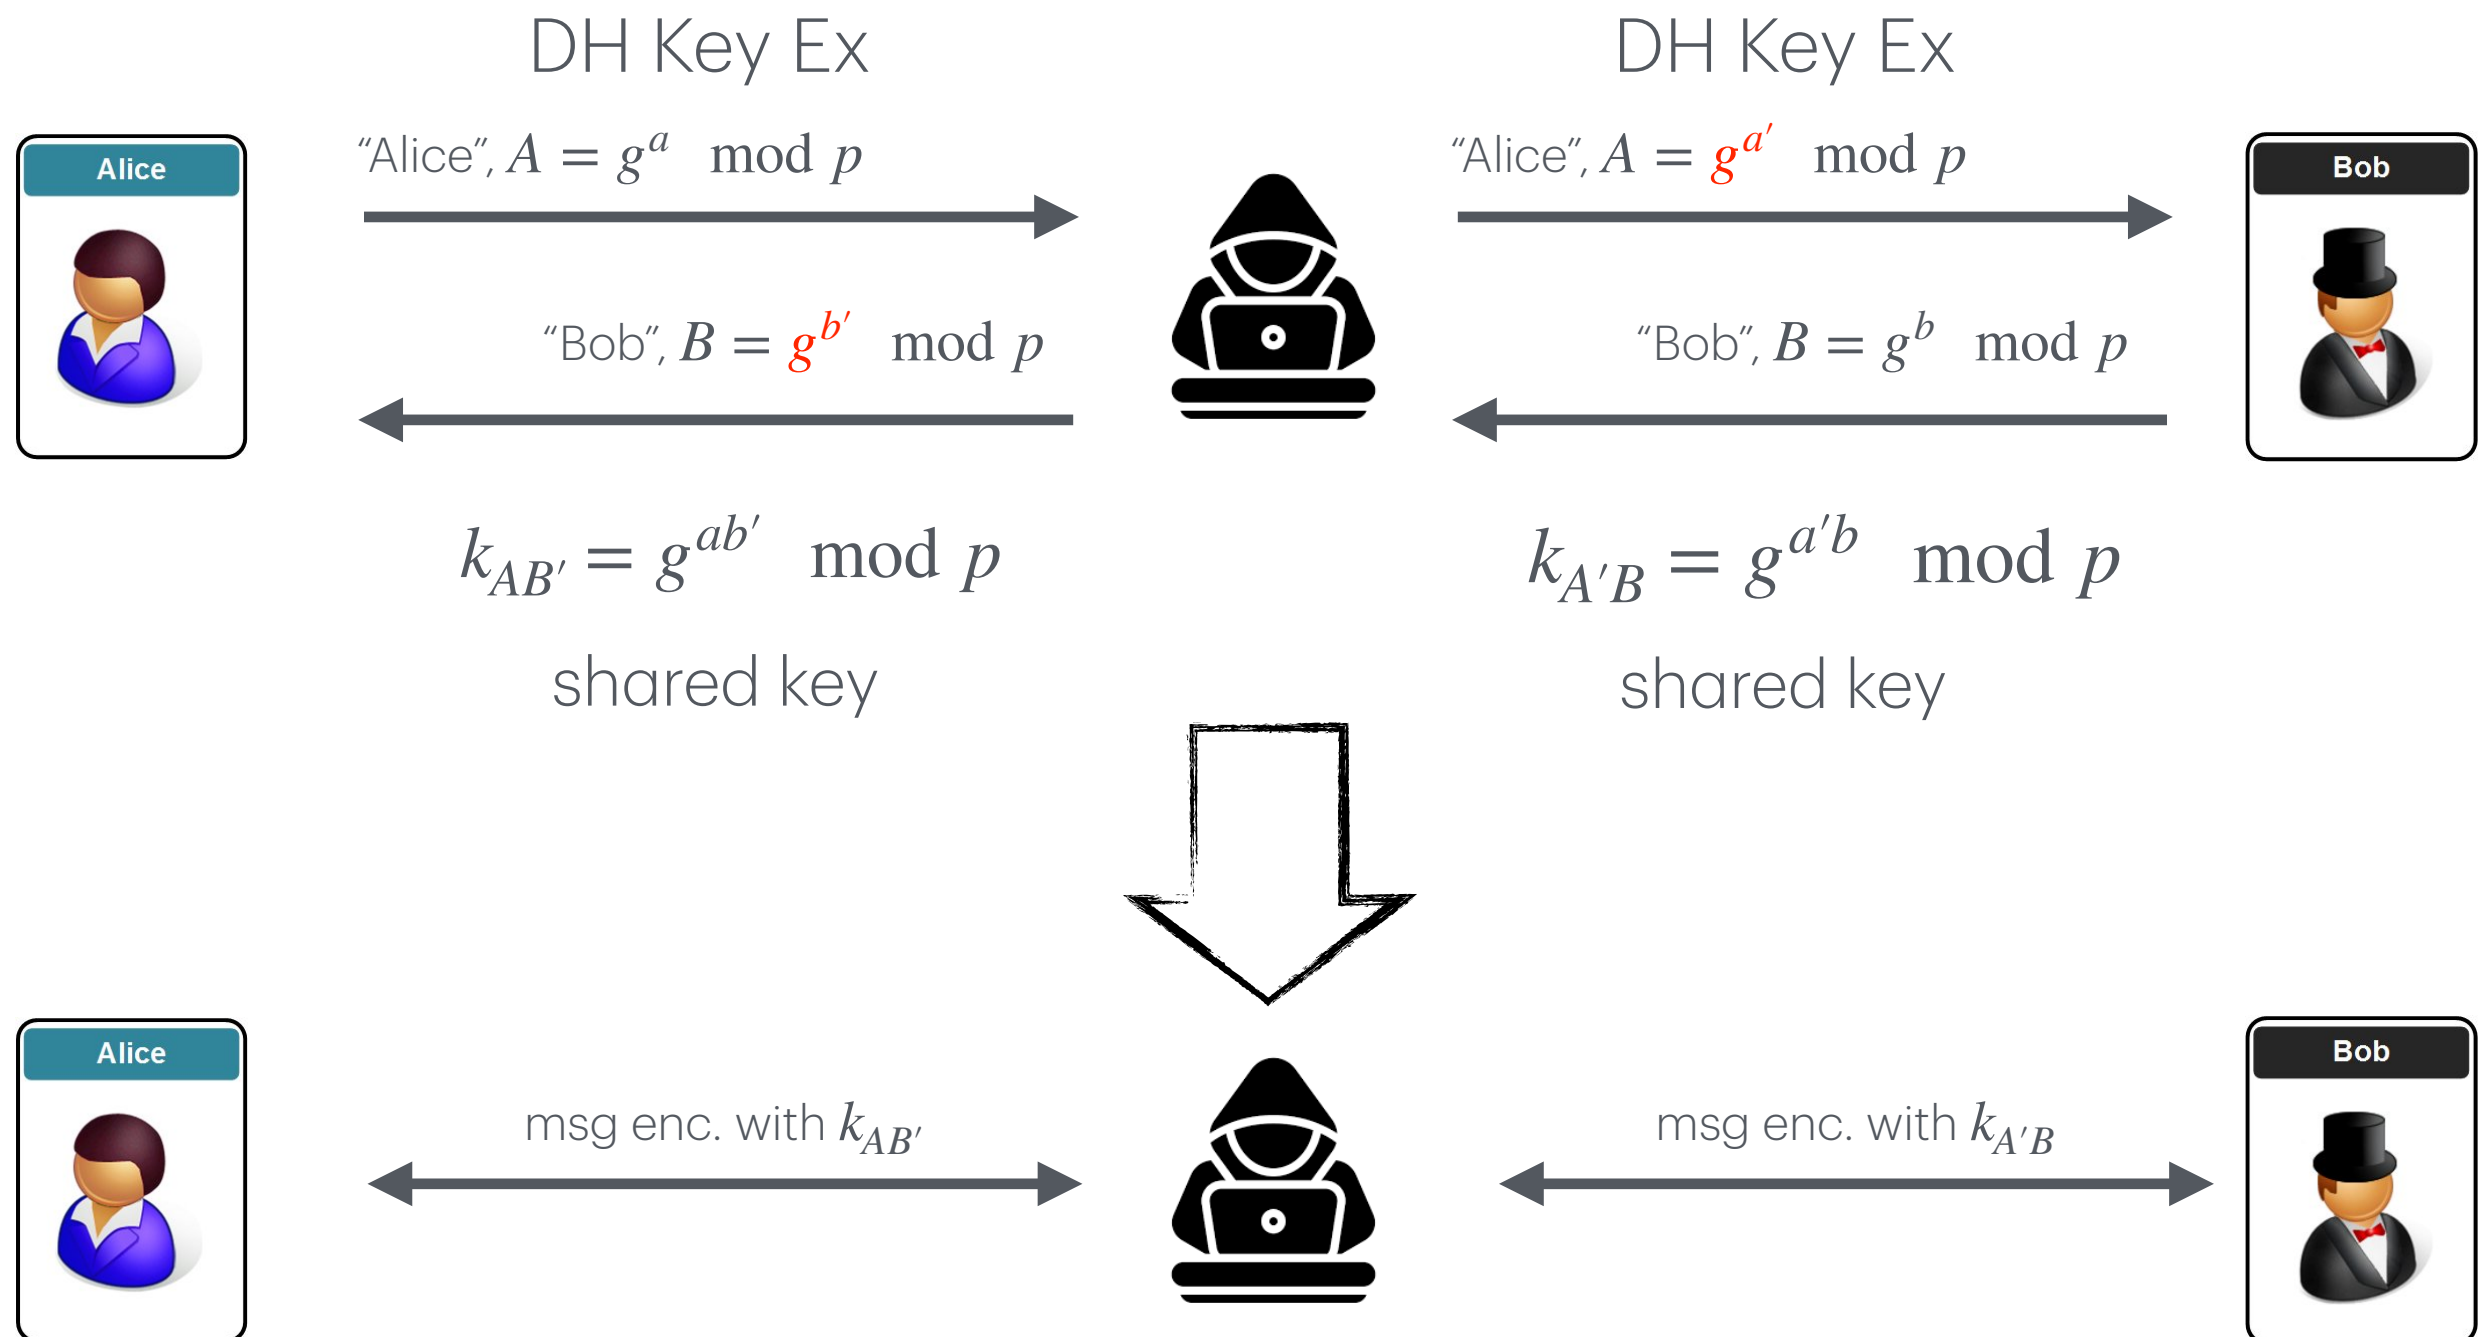
Security of Diffie-Hellman

- Suppose the prime p is n bits long.
 - Best known algorithm (GNFS): $\exp\left(\tilde{O}(\sqrt[3]{n})\right)$ run time

Cipher key size (e.g. AES)	Modulus size	Elliptic Curve size
80 bits	1024 bits	160 bits
128 bits	3072 bits	256 bits
256 bits	15360 bits	512 bits

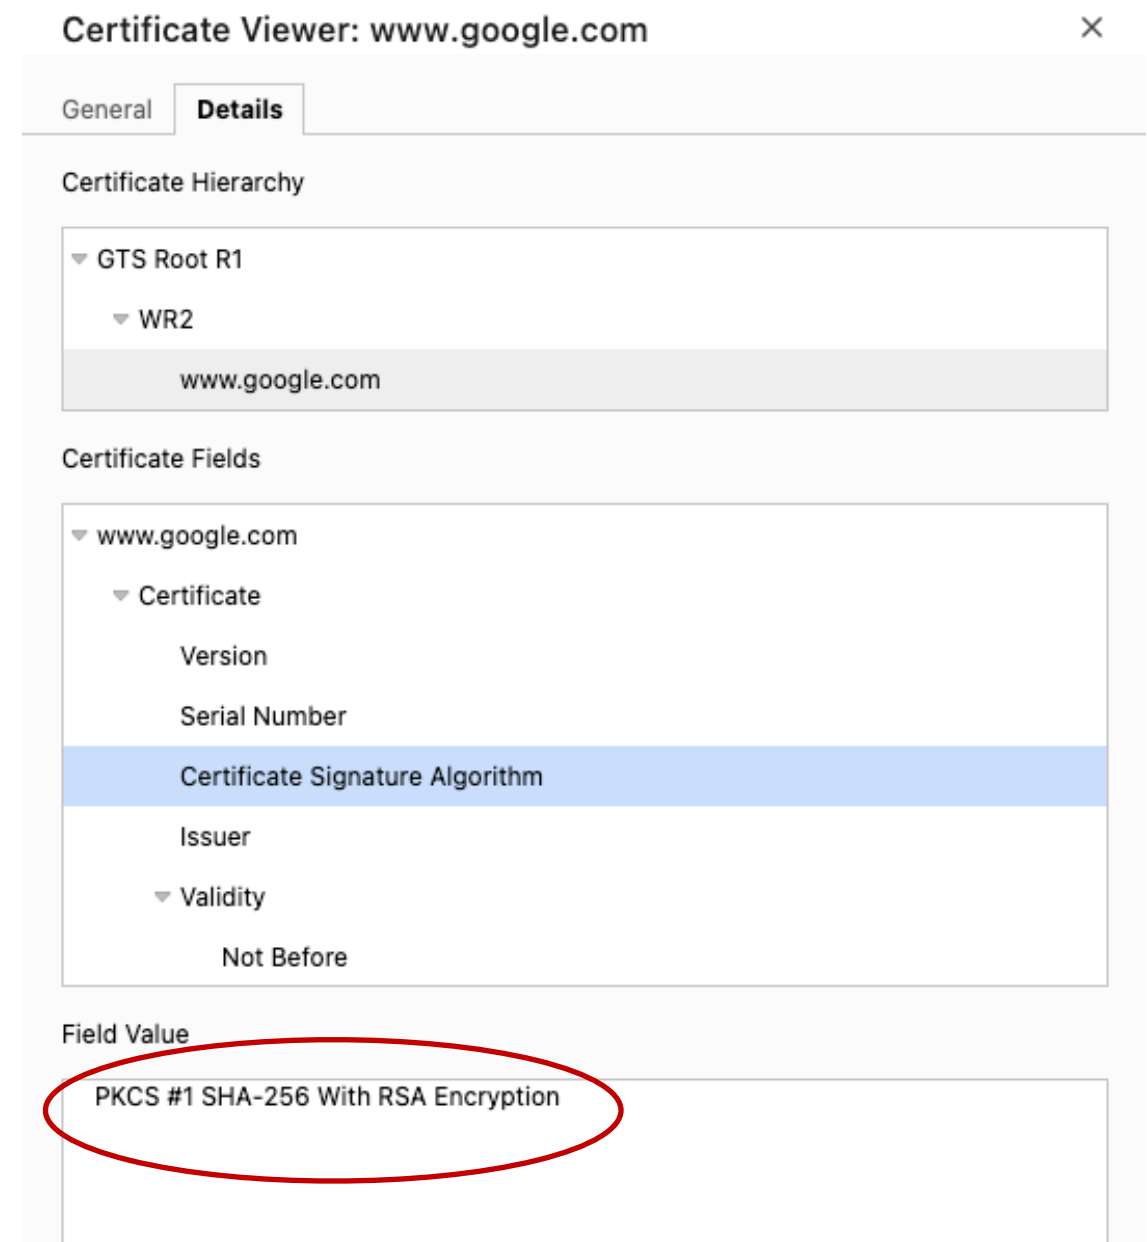
As a result: transition away from $\text{mod } p$ to elliptic curves

Insecure against Man-in-The-Middle



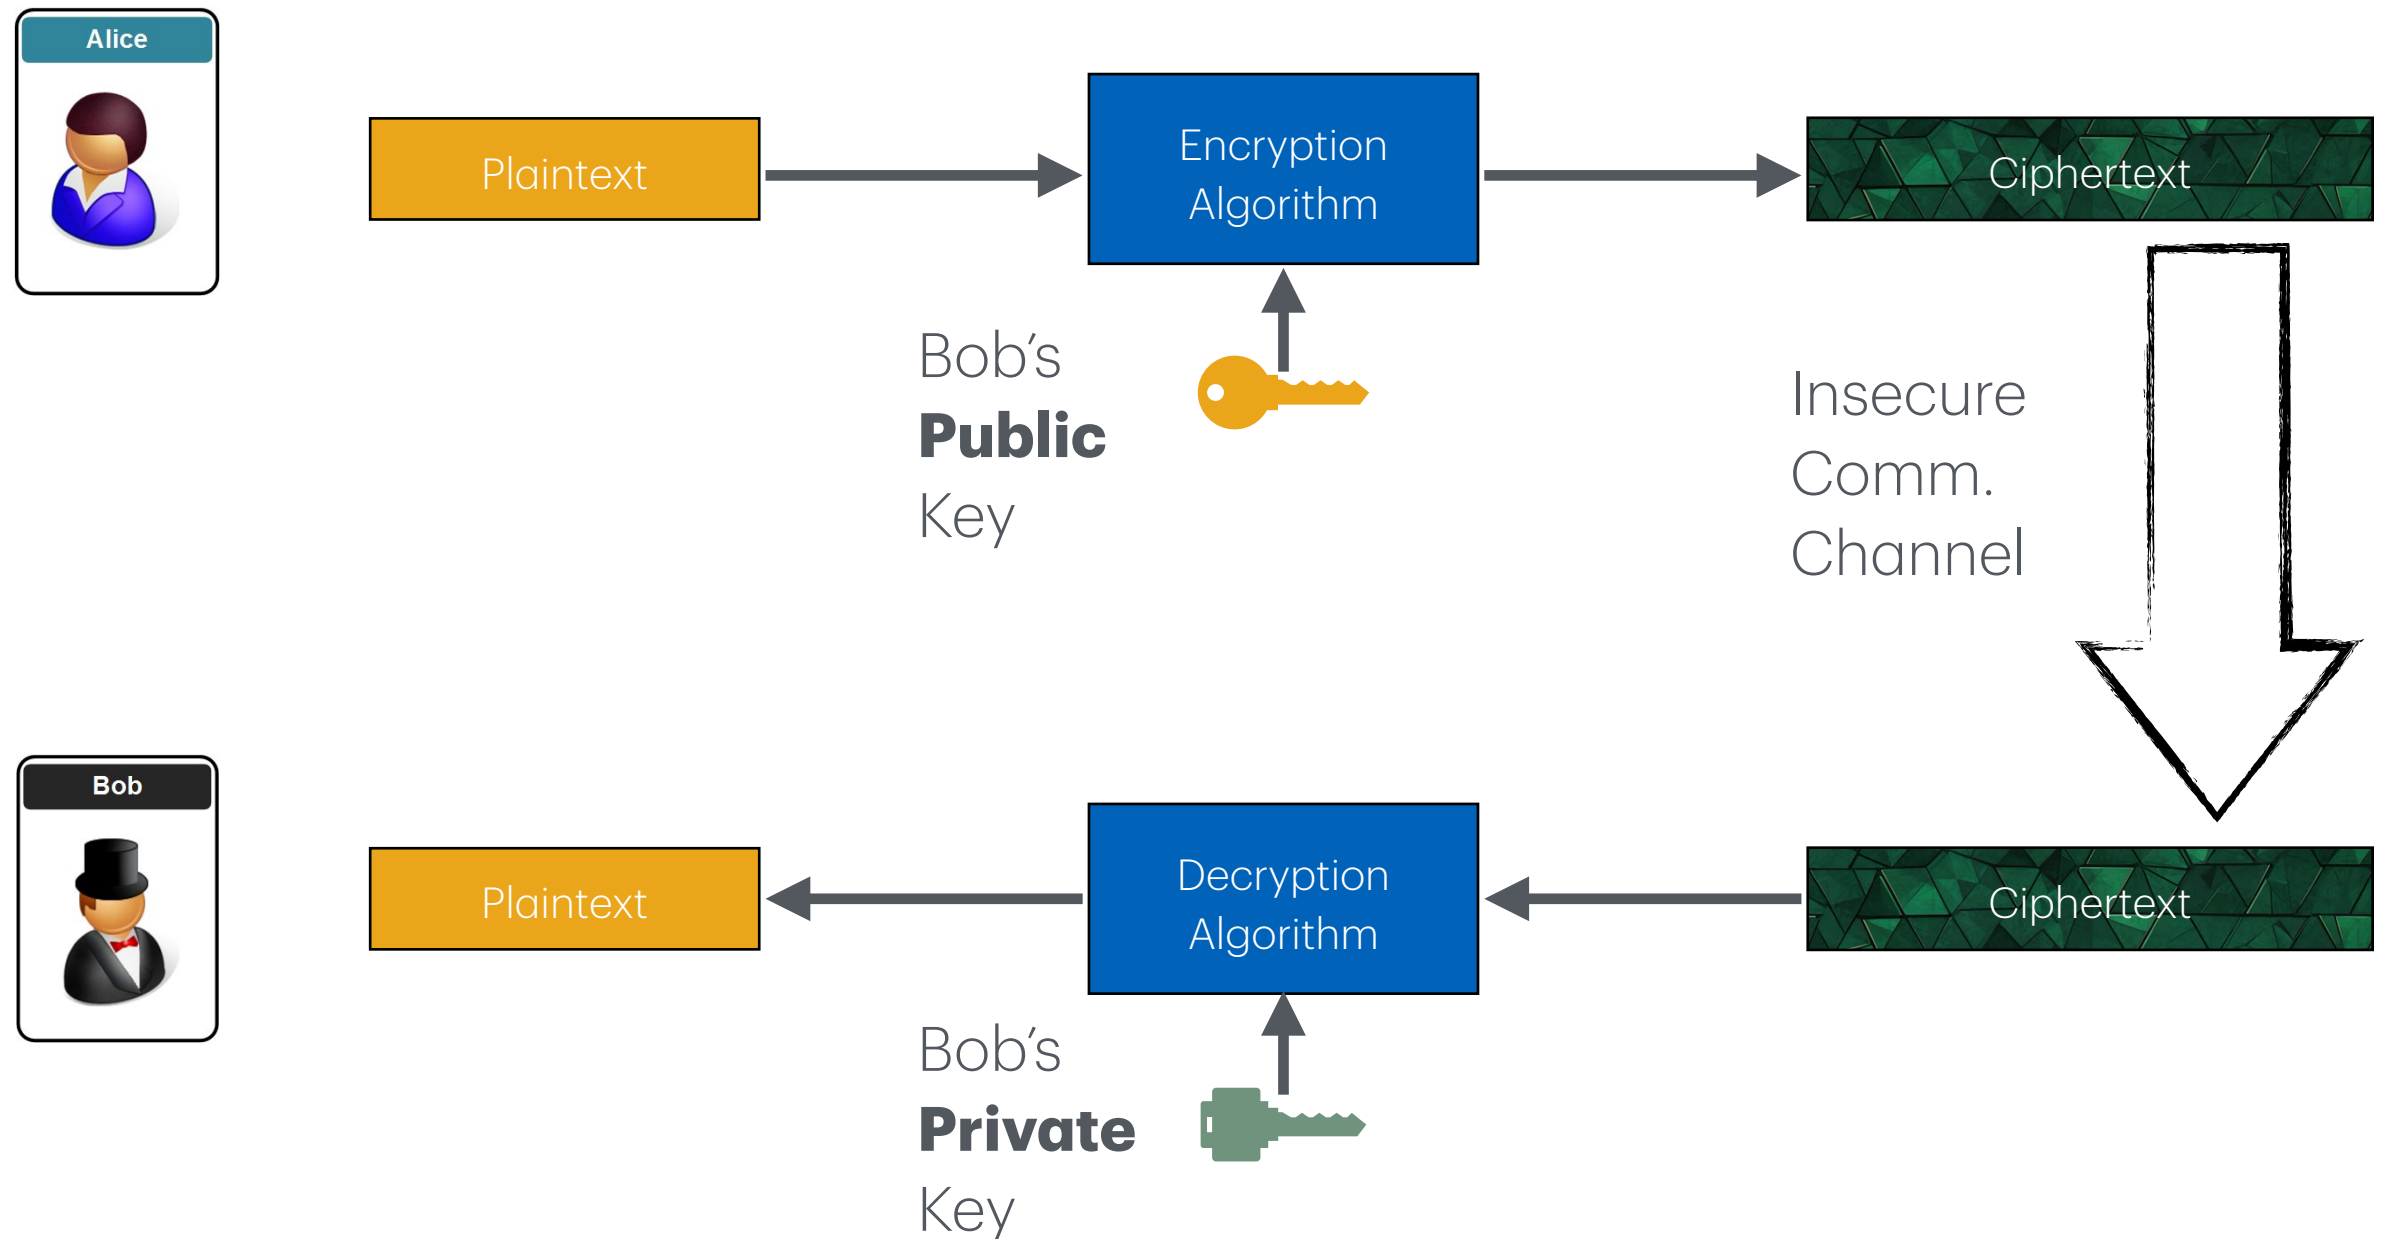
Insecure against Man-in-The-Middle

- Diffie-Hellman on itself does *not* provide authentication
- Always used along with *digital certificates* for authentication
- Based on PKC (usually RSA)
- Need a trusted third party (CA)



Public Key Cryptography (PKC)

Public Key Encryption



Public Key Encryption

- **Definition:** A triple of algorithms (G, E, D)
 - ▶ Key generation: *randomized* alg
 $G() \rightarrow (\text{public key } pk, \text{private key } sk)$
 - ▶ Enc.: *randomized* $E(pk, \text{message } m) \rightarrow \text{ciphertext } c$
 - ▶ Dec.: *deterministic* $D(sk, c) \rightarrow m$
- Consistency: $\forall (pk, sk)$ output by G :
 - ▶ $\forall \text{msg } m, D(sk, E(pk, m)) = m$

Application: Key Exchange


1. Generate key pair
 $(pk, sk) \leftarrow G()$



2. Send pk to Bob



3. generate symm
enc key $k \in K$



4. Send symm key enc
using $pk: c \leftarrow E(pk, k)$




5. decrypt to
get shared
symm key
 $k \leftarrow D(sk, c)$

PKC Security

- Like for Symm. Enc., we would require
 - **CPA-security**: Attacker provides a pair of eq. len. msgs (m_0, m_1) , and the server reply with $E(pk, m_b)$. Security means Attacker cannot infer b .
 - ▶ E must be randomized, otherwise this is trivially broken.
 - **CCA-security**:
 - A bit complicated
 - Recall in Symm. Enc., CCA-security is implied by Ciphertext Integrity.
 - ▶ I.e. Attacker cannot fabricate new valid ciphertext
 - But in PKC, anyone knowing pk can make new valid ciphertexts

PKC Security: CCA



PKC Security v.s. Symm Enc Security

- In symm-key settings:
 - **Ciphertext integrity**: Attacker cannot create *new* ciphertexts
 - Implies security against Chosen Ciphertext Attacks.
 - Implies authenticity.
- In public-key settings:
 - Attacker can create new ciphertexts using *pk* !!
 - So instead: we directly require Chosen Ciphertext Security.
 - But *no* authenticity: Still vulnerable to Man-in-The-Middle, like in the case of Diffie-Hellman.

PKC: Trapdoor Function

Trapdoor functions (TDF)

- A trapdoor function from $X \mapsto Y$ is the triple (G, F, F^{-1})
 - Key generation: *randomized* $G() \rightarrow$ (public key pk , private key sk)
 - $F(pk, \cdot)$: *deterministic* alg. that defines a function $X \mapsto Y$.
 - $F^{-1}(sk, \cdot)$: a function $Y \mapsto X$ that inverts $F(pk, \cdot)$
 - More specifically: $\forall (pk, sk)$ output by G , $\forall x \in X$,
 $F^{-1}(sk, F(pk, x)) = x$

Secure Trapdoor functions (TDF)

(G, F, F^{-1}) is **secure** if $F(pk, \cdot)$ cannot be inverted without sk

$(pk, sk) \leftarrow G()$

$x \sim X$



$(pk, y = F(pk, x))$



$x?$

The trapdoor is **secure** if any *computationally-bounded* attacker cannot infer x significantly better than random guess.

PKC Enc System from TDFs

Given

- (G, F, F^{-1}) : secure TDF $X \mapsto Y$
- (E_s, D_s) : symmetric auth. encryption defined over (K -key, M -msg, C -ciphertext)
- $H : X \mapsto Y$ a hash function

We construct a public-key enc. system(G, E, D):

- ▶ Key generation \mathbf{G} : same as G for TDF

PKC Enc System from TDFs

- Given:
 - Trapdoor function (G, F, F^{-1}) from $X \mapsto Y$
 - symmetric cipher $S = (E_S, D_S)$, hash function $H : Y \mapsto$ key space of S
- Key generation: $G() \longrightarrow$ public key pk , private key sk


Encryption $E(pk, m)$:

- Choose a random x from the key space of S
- Symm. secret key $h \leftarrow H(x)$
- Mask x as $y \leftarrow F(x)$
- Output $(y, E_S(h, m))$

Decryption $D(sk, (y, c))$:

- $x \leftarrow F^{-1}(y)$
- Symm. secret key $h \leftarrow H(x)$
- Output $D_S(h, c)$


PKC Enc System from TDFs



Theorem

- If (G, F, F^{-1}) is a secure TDF, (E_s, D_s) provides auth. enc.
- and $H : X \mapsto K$ is a “random oracle” (ideal hash func)
- then (G, E, D) is CCA secure (w.r.t. random oracle).

PKC Enc System from TDFs



Why so complicated?

Why not just send $c = F(m)$?

- Not CPA-secure because F is deterministic
- Why hash x as symm. enc. key?
 - $H(\cdot)$ destroys possible undesired structure in x
- Efficiency?
 - Can encrypt large message m outside F 's domain

PKC: RSA

Some basic number theory

Mod- n arithmetic

- $a \equiv_n b$ (or $a \equiv b \pmod n$) if the remainders of a and b divided by n are equal.
- Example: $n = 15$
 - $2 \times 8 = 16 = 15 + 1 \equiv_{15} 1$
 - $4 \times 8 = 32 = 2 \times 15 + 2 \equiv_{15} 2$
 - $6 \times 8 \equiv_{15} 3, 8 \times 8 \equiv_{15} 4, \dots$

Some basic number theory

Mod- n arithmetic

- If $ab \equiv_n 1$, then b is the (multiplicative) inverse of a (and vice versa)
- Example: $n = 15$
 - $2 \times 8 \equiv_{15} 1 \implies 2$ and 8 are inverse of each other
 - $3 \times ? \equiv_{15} 1 \implies 3$ does not have an inverse!
- Fact: Given n , one can find any a 's inverse (if exists) efficiently by the Extended Euclidean Algorithm

Some basic number theory

Mod- n arithmetic

- Let $\mathbb{Z}_n := \{0, 1, 2, \dots, n-1\}$, $\mathbb{Z}_n^* := \{x \in \mathbb{Z}_n : x \text{ is invertible}\}$

Fact: $x \in \mathbb{Z}_n^*$ if and only if $\gcd(x, n) = 1$

Theorem (Euler): For any two primes p, q , let $n := pq$, $\phi(n) := (p-1)(q-1)$, then $\forall x \in \mathbb{Z}_n^*, x^{\phi(n)} \equiv_n 1$

Corollary: For any two primes p, q , let $n := pq$, $\phi(n) := (p-1)(q-1)$, then $\forall x \in \mathbb{Z}_n, x^{\phi(n)} \cdot x \equiv_n x$

Some basic number theory

Mod- n arithmetic

- Let $\mathbb{Z}_n := \{0, 1, 2, \dots, n-1\}$, $\mathbb{Z}_n^* := \{x \in \mathbb{Z}_n : x \text{ is invertible}\}$
- Example: $n = 15 = 3 \times 5$, $\phi(n) = (3-1) \times (5-1) = 8$
 - Euler's Thm: recall $2 \times 8 \equiv 1 \pmod{15}$
 - $2^8 = 256 = 15 \times 15 + 1 \equiv 1 \pmod{15}$
 - $8^8 = 15 * 1118481 + 1 \equiv 1 \pmod{15}$
 - Corollary of Euler's Thm: recall 3 has no multiplicative inverse module 15
 - $3^8 = 437 \times 15 + 6 \not\equiv 1 \pmod{15}$, but $3^{8+1} = 1312 \times 15 + 3 \equiv 3 \pmod{15}$

The RSA Trapdoor Perm.

$(G, \text{RSA}, \text{RSA}^{-1})$

- Intuition: want pub-key e and priv-key d s.t.
 - $\text{RSA}(m) = m^e \pmod N$
 - $\text{RSA}^{-1}(c) = c^d = m^{ed} = m \pmod N$

The RSA Trapdoor Perm.

$(G, \text{RSA}, \text{RSA}^{-1})$

- Intuition: want pub-key e and priv-key d s.t.

- $\text{RSA}(m) = m^e \bmod N$

Coro of Euler's Thm: $ed = \alpha\phi(N) + 1$

- $\text{RSA}^{-1}(c) = c^d = m^{ed} = m \bmod N$

- RSA Key generation $G()$:

- pick two random primes $p, q \approx 1024$ bits, set $N = pq$.
- Pick two integers e, d s.t. $ed \equiv_{\phi(N)} 1$
- Output Pub-key $pk = (N, e)$, private key $sk = (N, d)$

The RSA Trapdoor Perm.

$(G, \text{RSA}, \text{RSA}^{-1})$

- Why is it a *trapdoor*?
- To invert $\text{RSA}(\cdot)$ (without priv-key d), need to solve x from $c = x^e \pmod N$

Best known alg.

- Factor N to get p, q
- $\phi(N) \leftarrow (p - 1)(q - 1)$
- Find d as e 's inverse (mod $\phi(N)$) using Ext. Euclidean Alg
- $x \leftarrow c^d \pmod N$

The RSA Trapdoor Perm.

$(G, \text{RSA}, \text{RSA}^{-1})$

- Why is it a *trapdoor*?
- To invert $\text{RSA}(\cdot)$ (without priv-key d), need to solve x from $c = x^e \bmod N$ (discrete root)

Best known alg.

- Factor N to get $p, q \longleftarrow$ believed to be hard for large N
- $\phi(N) \leftarrow (p - 1)(q - 1)$
- Find d as e 's inverse (mod $\phi(N)$) using Ext. Euclidean Alg
- $x \leftarrow c^d \bmod N$

Is there a / no shortcut without factoring N ?

Ans: unknown.

The RSA Public-key Enc. (ISO)

- Given:
 - symmetric cipher $S = (E_S, D_S)$,
 - hash function $H : \mathbb{Z}_N \mapsto$ key space of S
- Key generation $G()$: The RSA trapdoor key generation
 - Output pub-key $pk = (N, e)$, private key $sk = (N, d)$

Encryption $E(pk, m)$:

- Choose random x from \mathbb{Z}_N
- $y \leftarrow \text{RSA}(x) = x^e \bmod N$
- symm. enc. key $h \leftarrow H(x)$
- Output $(y, E_S(h, m))$

Decryption $D(sk, (y, c))$:

- $x \leftarrow \text{RSA}^{-1}(y)$
- Symm enc key $h \leftarrow H(x)$
- Output $D_S(h, c)$

PKC: RSA in practice

Rule of Thumb

- Never use the RSA trapdoor function directly as encryption scheme
 - Easily breakable by CCA (how?)
- Do not implement the alg by yourself.
 - So many ways to go wrong.
 - Use well-established libraries, e.g. OpenSSL, libsodium, etc.

Parameter choices

- N should be at least 2048 bits, better 4096 bits
 - $p, q \sim$ at least 1024 bits
- Use small e to speed up encryption: $e = 3, 5, 65537$
- Never use a small d :
 - Very small d : attacker can enumerate it
 - Moderately small: $d < N^{0.292}$ is still insecure [Boneh-Durfee'98]

Padding

- Recall the RSA PKC encryption:

Encryption $E(pk, m)$:

- Choose random x from \mathbb{Z}_N
- $y \leftarrow \text{RSA}(x) = x^e \bmod N$
- symm. enc. key $h \leftarrow H(x)$
- Output $(y, E_S(h, m))$

- We *first* sample a x from the RSA TDF domain (e.g. 2048 bits int when $|N| = 2048$).
- Then the key h for symm enc is obtained by hashing x

Padding

- In practice the order is often reversed:
 - ▶ We first generate the key h for symm. enc. (e.g., 128 bits for AES-128)
 - ▶ Then h is masked using TDF and sent along with encrypted msg.

Encryption $E(pk, m)$:

- Choose random symm enc key
- $y \leftarrow \text{RSA}(h) = h^e \bmod N$
- Output $(y, E_S(h, m))$

Problem: h is usually too smaller than N (e.g., 128 bits vs 2048 bits)


Padding

- Example:
 - Suppose $e = 3$, then for a small h , it's likely that $h^e < N$
 - The attacker can get h by computing $y^{1/3}$ (*without mod N* , so this is just a regular cubic root calculation)

Padding

- Solution: **Pad** some bits to h to make it 2048 bits long

PKCS1 v1.5 Padding scheme:



- Widely deployed, e.g. in HTTPS - TLS 1.2
- Unfortunately, weakness known [Bleichenbacher'1998].
- Use RSA-OAEP instead.

PKC: ElGamal

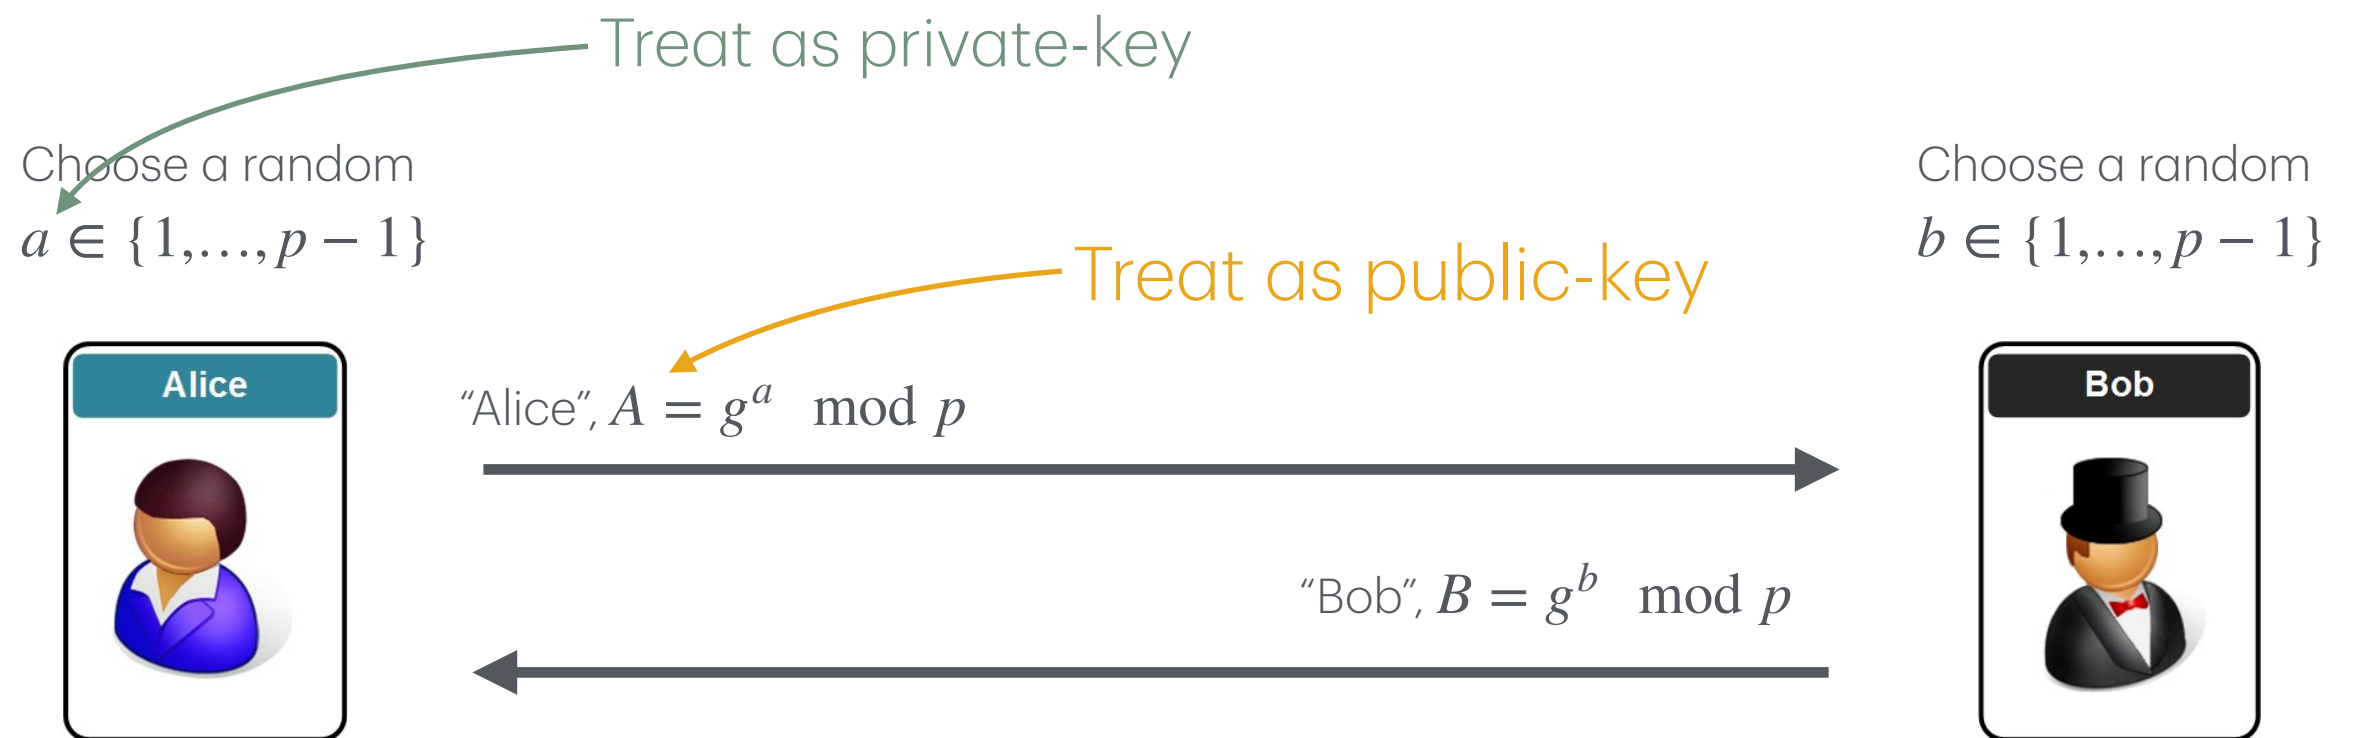
Other Public-Key Algorithms

- Many popular public key algorithms rely on the difficulty of *discrete logarithm problem*
 - ElGamal encryption and ElGamal signature
 - Digital Signature Algorithm (DSA)
 - Diffie-Hellman key exchange
 - ...
- Given an appropriate setup with g, p and $h = g^x \bmod p$, it is difficult for someone to compute x
 - ▶ x is called the discrete logarithm of h to the base g
 - ▶ groups in which the discrete logarithm problem is hard use prime modulus p (conventional and elliptic curve settings)

PKC from Discrete Log

Basic ideas

Recall Diffie-Hellman:



$$(g^b)^a = B^a \mod p = k_{AB} = g^{ab} \mod p = (g^a)^b = A^b \mod p$$

shared key

PKC from Discrete Log

Basic ideas

(E_s, D_s) : some symm. auth. enc.

Choose a random
 $sk = a \in \{1, \dots, p-1\}$

Choose a random
 $b \in \{1, \dots, p-1\}$


$$pk = A = g^a \bmod p$$



Compute $g^{ab} = A^b$;

Derive k from g^{ab}

$$[B = g^b, c = E_s(k, m)]$$



Compute $g^{ab} = B^a$;

Derive k from g^{ab} ;

$$m \leftarrow D_s(k, c)$$

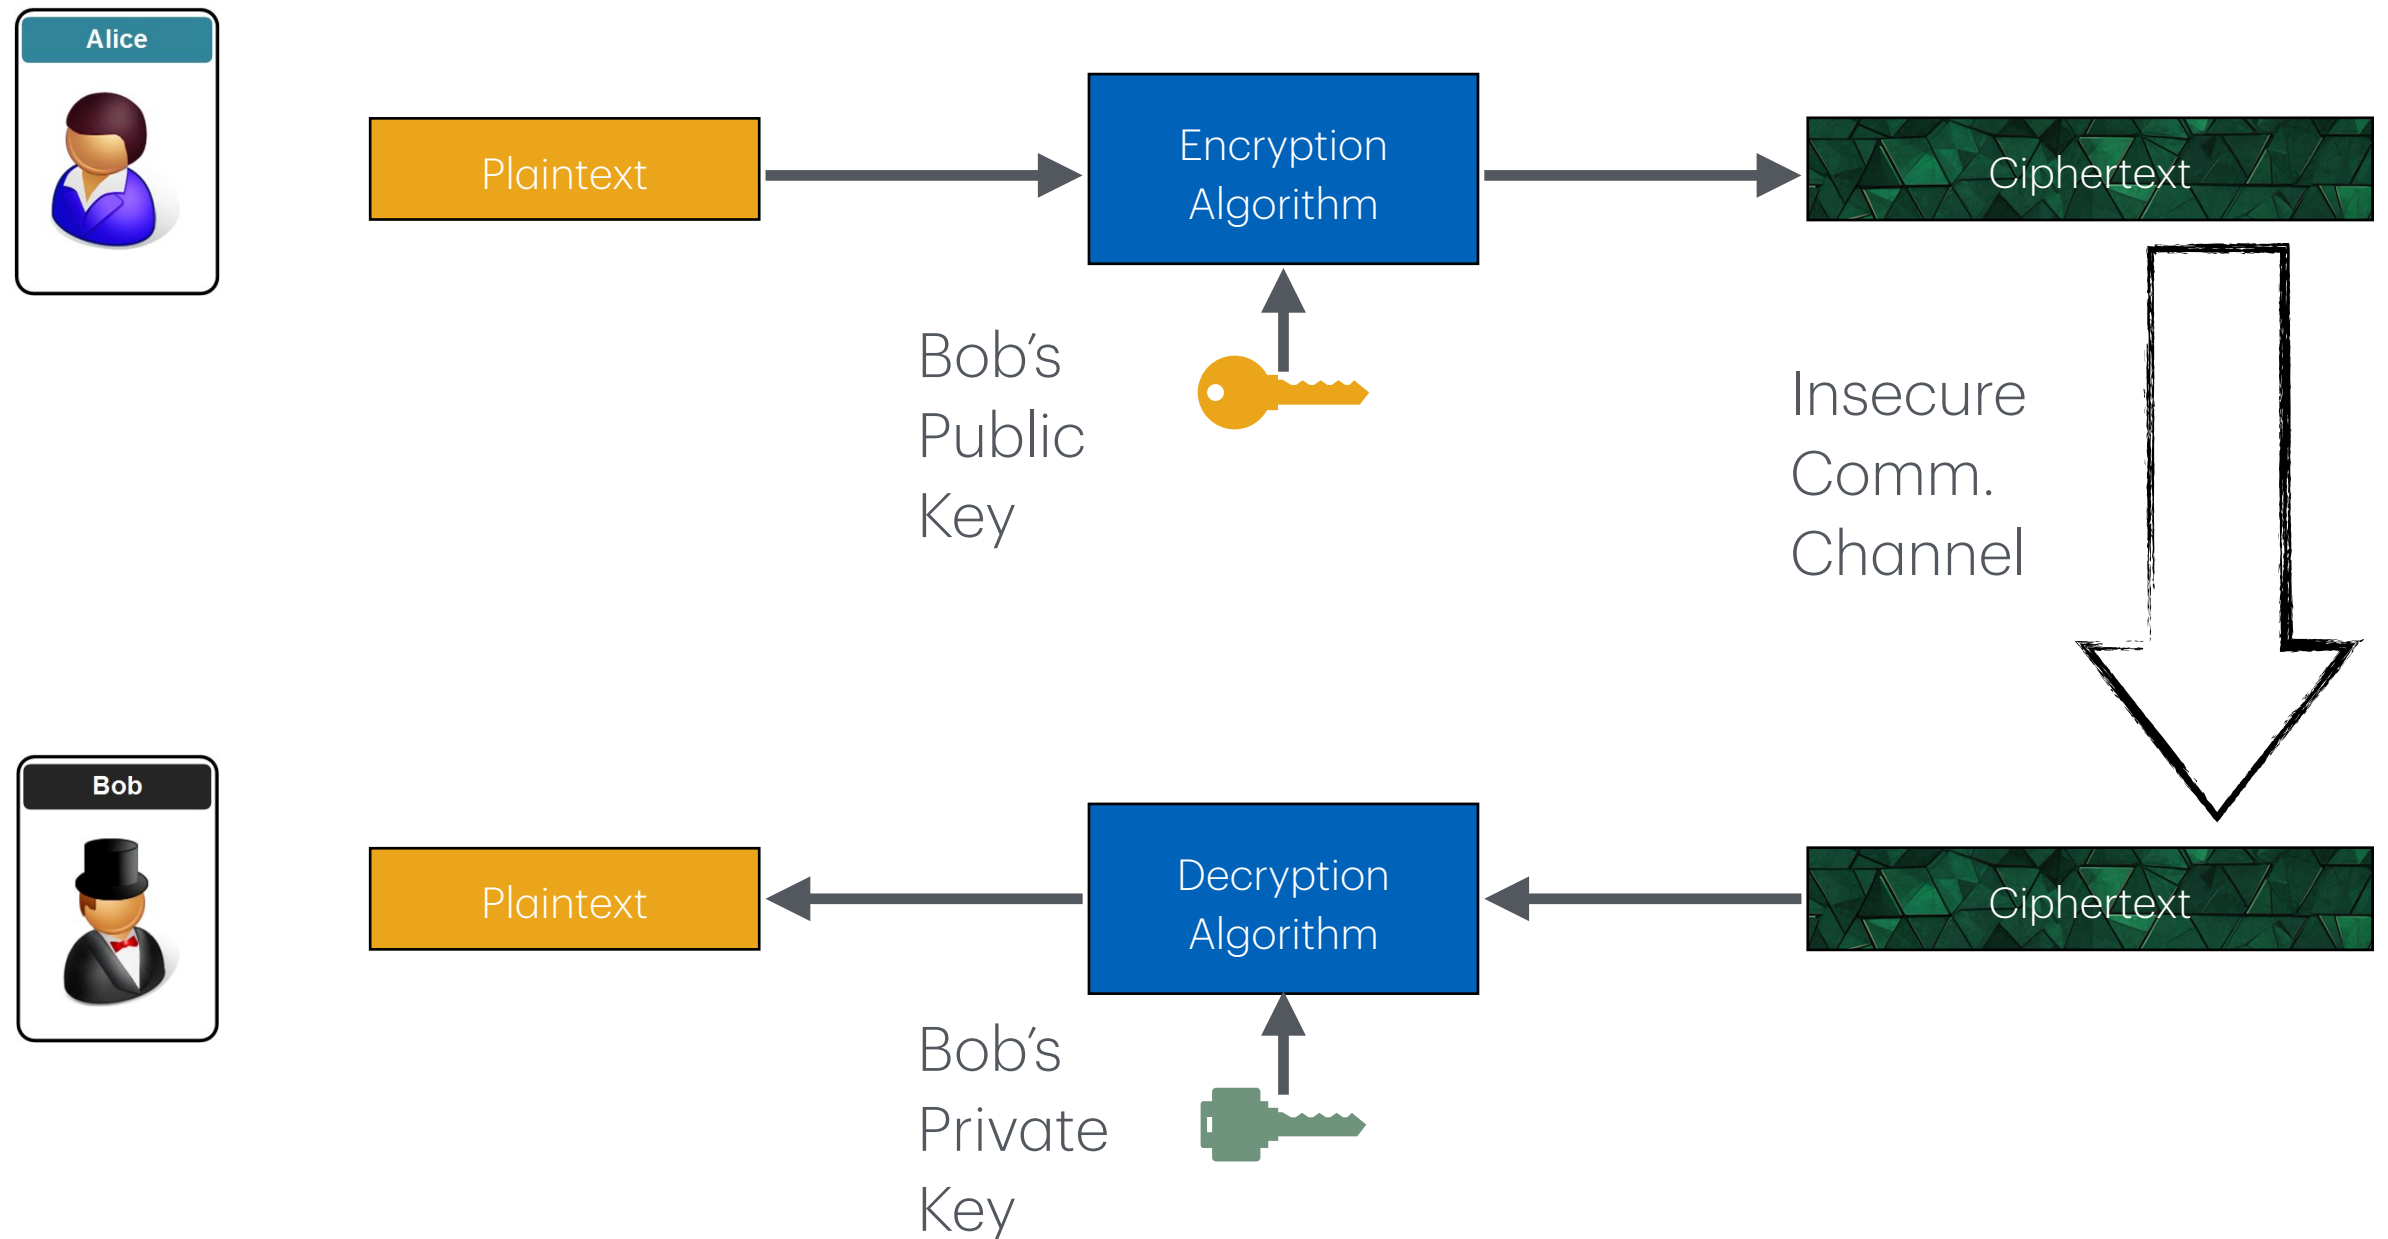
In actual impl. g is also
chosen randomly and
published as part of the pk

Symmetric vs Public-Key Encryption

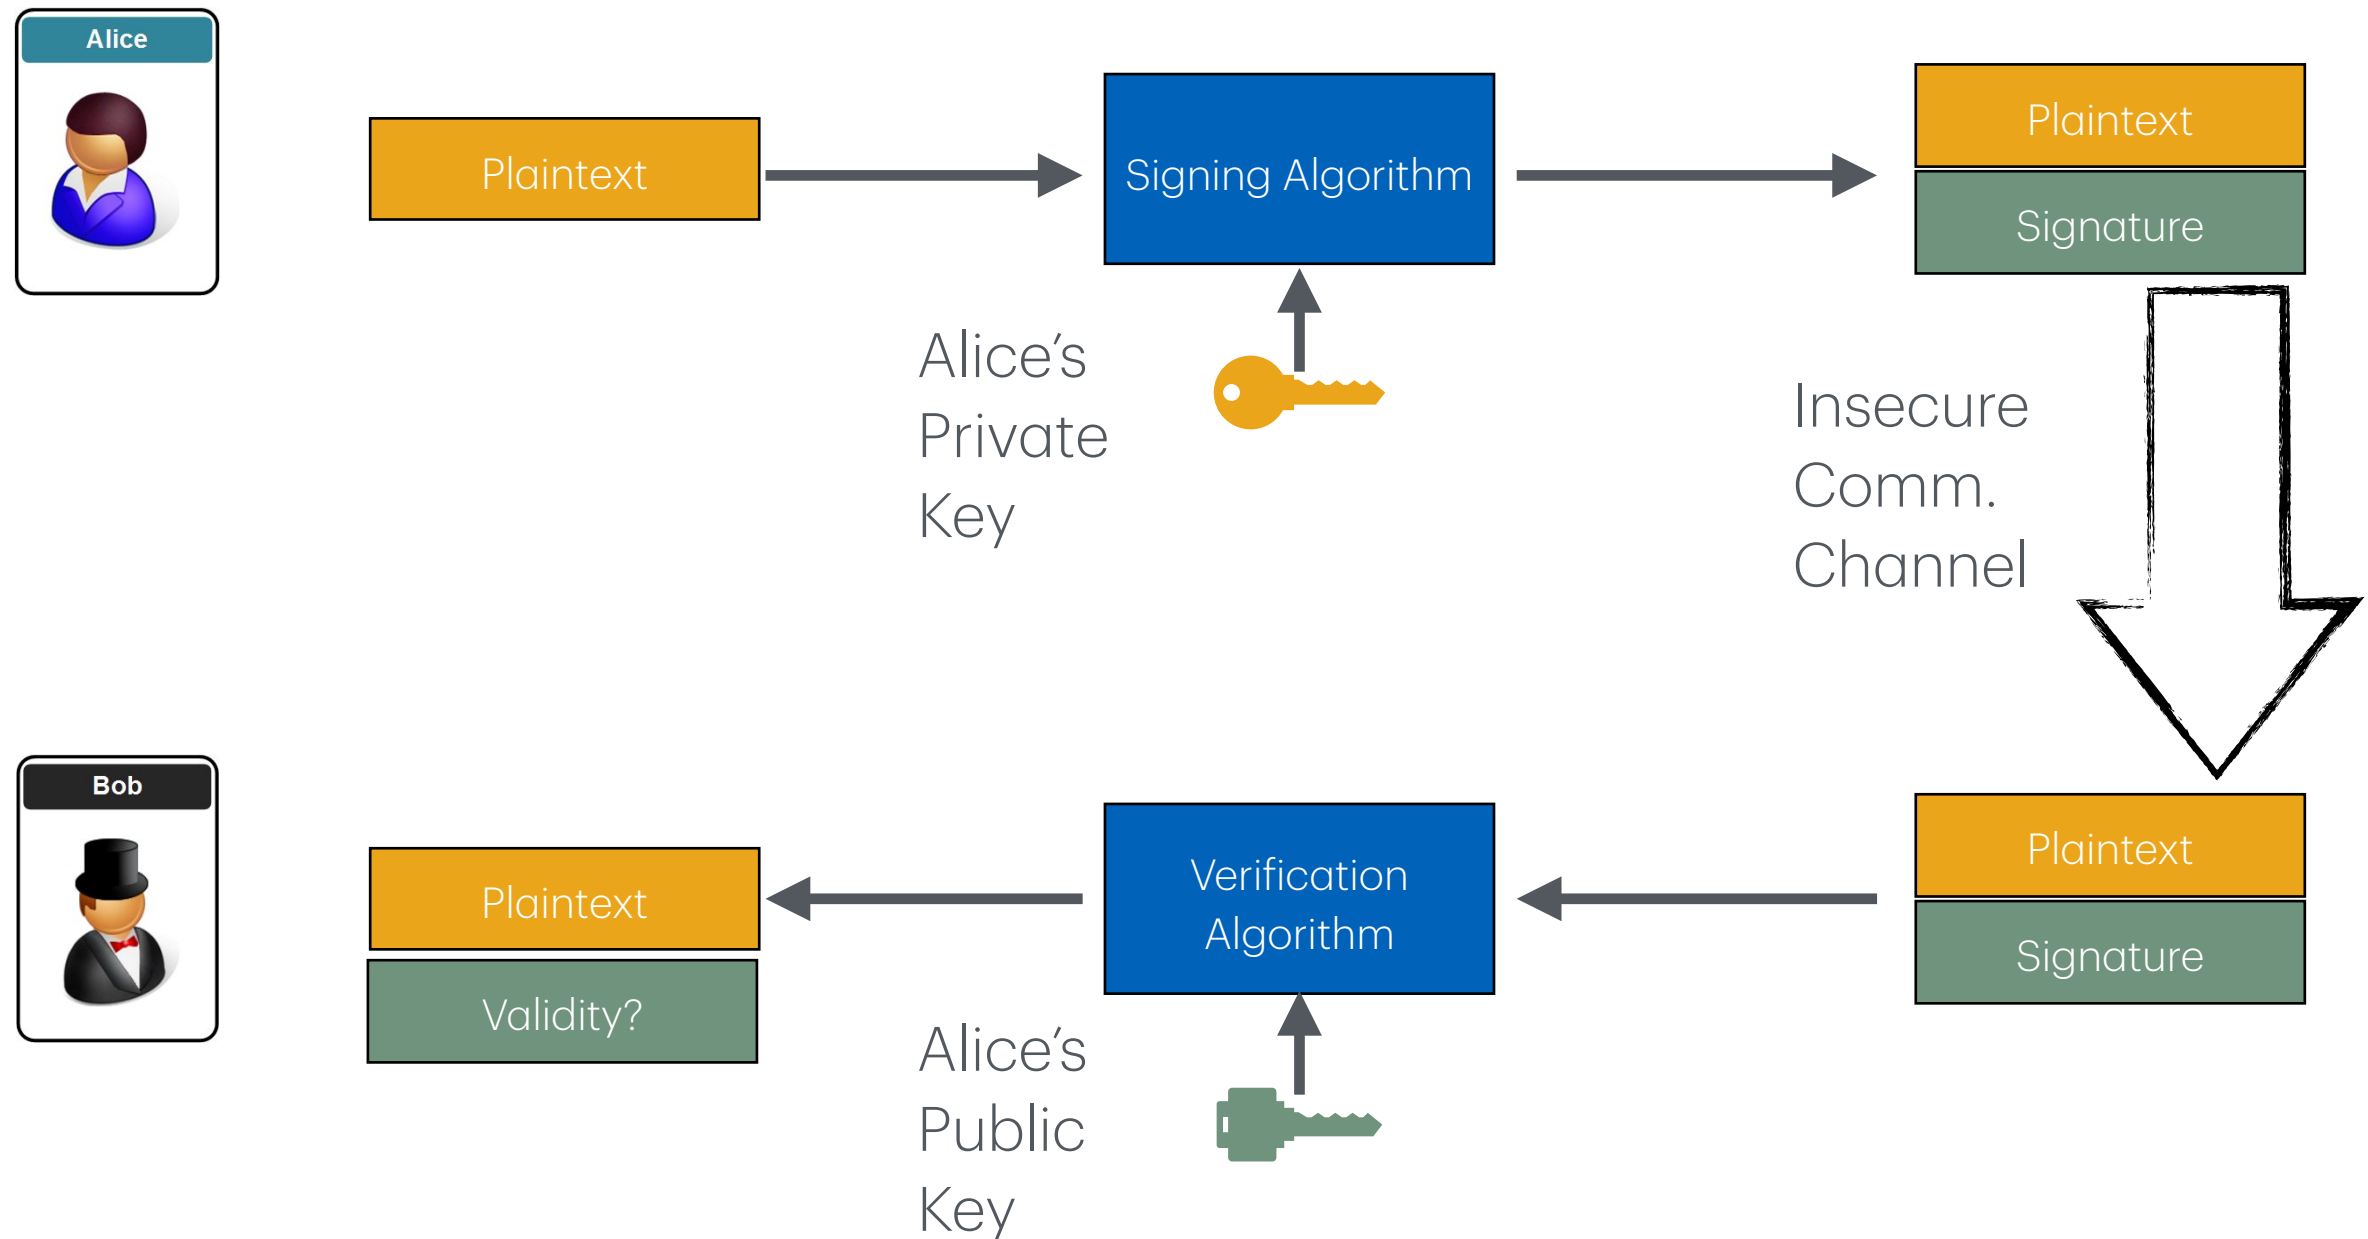
- Public-key operations are orders of magnitude slower than symmetric encryption
 - A multiplication modulo n requires close to $O(\log^2 n)$ work
 - A full-size exponentiation modulo n requires close to $O(\log^3 n)$ work
 - It is the cost of multiplication times the exponent size
 - Public-key encryption is typically not used to communicate large volumes of data
 - It is rather used to communicate (or agree on) a symmetric key
 - The data itself is sent encrypted with the symmetric key
- In RSA, decryption is significantly slower than encryption, with key generation being the slowest

PKC: Digital Signature

Public-Key Encryption




(Public-Key) Digital Signature




Example: Digital Signature from TDF

Given: a secure trapdoor function (G, F, F^{-1}) and $(pk, sk) \leftarrow G()$

Signing alg. $S(sk, \text{msg})$



Verify alg. $V(pk, \text{msg}, \text{sig})$



Example: Digital Signature from RSA TDF

- Given:
 - RSA trapdoor function $(G, \text{RSA}, \text{RSA}^{-1})$,
 - collision-resistant hash function $H : \mathbb{Z}_N \mapsto \text{key space of } S$
- Key generation $G()$: The RSA trapdoor key generation
 - Output pub-key $pk = (N, e)$, private key $sk = (N, d)$

Signing $S(sk, m)$:

$$- y \leftarrow \text{RSA}^{-1}(H(m)) = H(m)^d \bmod N$$

In practice $H(m)$ will be padded before usage.

Verifying $V(pk, m, \sigma)$:

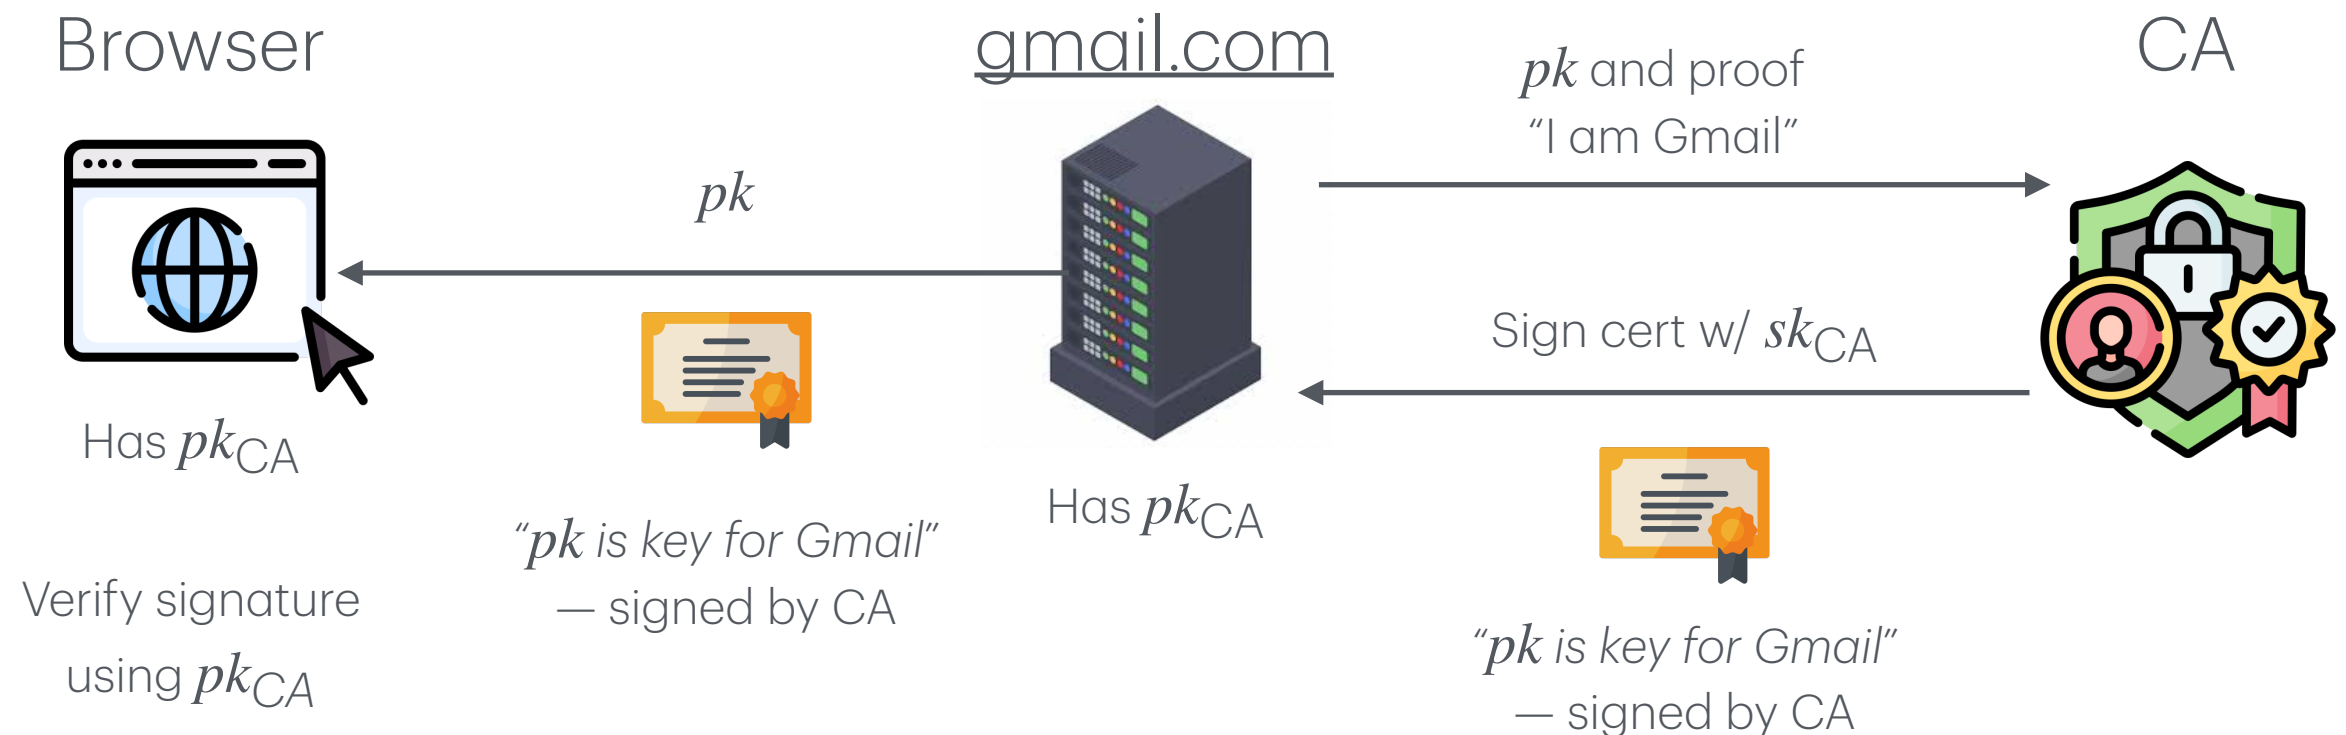
$$- \text{Check if } H(m) = \text{RSA}(\sigma) = \sigma^e \bmod N$$

When to use signatures

- If **one** party signs and **one** party verifies: use a *MAC*
 - ▶ Often requires interaction to generate a shared key
 - ▶ Recipient can modify the data and re-sign it before passing the data to a 3rd party
- If **one** party signs and **many** parties verify: use a *signature*
 - ▶ Recipients cannot modify received data before passing data to a 3rd party (non-repudiation)

Application: Certificate

- Problem: browser needs server's public-key to setup a session key
- Solution: server asks trusted 3rd party (CA) to sign its public-key pk .



The Big Picture

Security Goal	Symmetric Key		Public Key
Secrecy / Confidentiality	Block ciphers with encryption modes (AES); Stream ciphers	Authenticated Encryption	Public key encryption (RSA, ElGamal, etc)
Authenticity / Integrity	Message Authentication Codes (CBC-MAC, HMAC)		Digital signatures (RSA, DSA, etc)

Acknowledgement

- The slides of this lecture is developed heavily based on
 - Slides from Prof Dan Boneh's lecture on Cryptography (<https://crypto.stanford.edu/~dabo/courses/OnlineCrypto/>)
 - Slides from Prof Ziming Zhao's lecture on Computer Security (<https://zzm7000.github.io/teaching/2023springcse410565/index.html>)

Questions?