

Software Security I

CSE 565: Fall 2024
Computer Security

Xiangyu Guo (xiangyug@buffalo.edu)

University at Buffalo

Disclaimer

- We don't claim any originality of the slides. The content is developed heavily based on
 - Slides from ASU security team's pwn.college
 - Slides from Prof Ziming Zhao's past offering of CSE565 (<https://zzm7000.github.io/teaching/2023springcse410565/index.html>)
 - Slides from Prof Hongxin Hu's past offering of CSE565

Announcement

- HW3 and Project3 **due Tue, Nov 12, 23:59 pm.**
- Bonus in-class quizzes starting from next week

Review of Last Lecture

- Privacy
 - 3PD tracking
 - 3PD Cookies
 - Fingerprinting
- Anonymity
 - Tor network
- Secure messaging app
 - E2E Encryption
 - Deniability

Today's topic

- Background
 - ELF
- Linux process Loading & Execution

From a C program to a process


cat.c



```
// implements a cat!
int main(int argc, char **argv) {
    char buf[2048];
    int n;
    int fd = argc == 1 ? 0: open(argv[1], 0);
    while ((n = read(fd, buf, 2048)) > 0 && write(1, buf, n) > 0);
}
```

```
$ cat cat.c
// implements a cat!
int main(int argc, char **argv) {
    char buf[2048];
    int n;
    int fd = argc == 1 ? 0: open(argv[1], 0);
    while ((n = read(fd, buf, 2048)) > 0 && write(1, buf, n) > 0);
}
```

From a C program to a process



Pre-processing

```
seed@seed-vm ~/Programs> gcc -E cat.c -o cat.i
seed@seed-vm ~/Programs> cat cat.i
# 0 "cat.c"
# 0 "<built-in>"
# 0 "<command-line>"
# 1 "/usr/include/stdc-predef.h" 1 3 4
# 0 "<command-line>" 2
# 1 "cat.c"

int main(int argc, char **argv) {
    char buf[2048];
    int n;
    int fd = argc == 1 ? 0: open(argv[1], 0);

    while ((n = read(fd, buf, 2048)) > 0 && write(1, buf, n) > 0);
}
```

- No header/macro expansion because there was none in the source code.
- Comment is removed.

Compilation

```
seed@seed-vm ~/Programs> gcc -S cat.i -o cat.s -Wno-implicit-function-declaration
seed@seed-vm ~/Programs> head -40 cat.s
        .arch armv8-a
        .file     "cat.c"
        .text
        .align    2
        .global   main
        .type     main, %function
main:
.LFB0:
        .cfi_startproc
        sub      sp, sp, #2096
        .cfi_def_cfa_offset 2096
        stp      x29, x30, [sp]
        .cfi_offset 29, -2096
        .cfi_offset 30, -2088
        mov      x29, sp
        str      w0, [sp, 28]
        str      x1, [sp, 16]
        adrp     x0, :got:__stack_chk_guard
        ldr      x0, [x0, #:got_lo12:__stack_chk_guard]
        ldr      x1, [x0]
        str      x1, [sp, 2088]
```

Assembly

```
seed@seed-vm ~/Programs> gcc -c cat.s -o cat.o
```

```
seed@seed-vm ~/Programs> strings cat.o
```

```
GCC: (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0
```

```
cat.c
```

```
main
```

```
__stack_chk_guard
```

```
open
```

```
read
```

```
write
```

```
__stack_chk_fail
```

```
.symtab
```

```
.strtab
```

```
.shstrtab
```

```
.rela.text
```

```
.data
```

```
.bss
```

```
.comment
```

```
.note.GNU-stack
```

```
.rela.eh_frame
```

```
seed@seed-vm ~/Programs> cat cat.o
```

```
ELFx@@
```

```
{
```

```
@@@qT
```

```
T'@ Rq,TR*@D@c@T*{@ _GCC: (Ubuntu 11.4.0-1ubuntu1~22.04) 11.4.0zRx
```

```
AAq
```

```
$).4cat.c$x$dmain__stack_chk_guardopenreadwrite__stack_chk_fail7
```

```
8
```

```
D
```

```
8
```

```
yntab.strtab.shstrtab.rela.text.data.bss.comment.note.GNU-stack.rela.eh_frame @@ &10, :<0@J@
```

```
x
```

```
EY<
```

Linking

```
seed@seed-vm ~/Programs> gcc -o cat cat.o
seed@seed-vm ~/Programs> file cat.o
cat.o: ELF 64-bit LSB relocatable, ARM aarch64, version 1 (SYSV), not stripped
seed@seed-vm ~/Programs> file cat
cat: ELF 64-bit LSB pie executable, ARM aarch64, version 1 (SYSV), dynamically linked,
interpreter /lib/ld-linux-aarch64.so.1, BuildID[sha1]=e1dabc44b4317a3354cf92e32db4a442f
68b20f6, for GNU/Linux 3.7.0, not stripped
```

ELF Binary Format

What is a Cat

```
> file ./cat
```

```
./cat: ELF 64-bit LSB pie executable, ARM aarch64, version 1 (SYSV),  
dynamically linked, interpreter /lib/ld-linux-aarch64.so.1, for GNU/  
Linux 3.7.0, BuildID[sha1]=e1dabc44b4317a3354cf92e32db4a442f68b20f6,  
not stripped
```

What is a Cat

```
> file ./cat
```

```
./cat: ELF 64-bit LSB pie executable, ARM aarch64, version 1 (SYSV),  
dynamically linked, interpreter /lib/ld-linux-aarch64.so.1, for GNU/  
Linux 3.7.0, BuildID[sha1]=e1dabc44b4317a3354cf92e32db4a442f68b20f6,  
not stripped
```

What is an ELF?

- **ELF** (Executable and Linkable Format) is a [binary](#) file format.
 - Contains the program and its data.
 - Describes how the program should be loaded (*program/segment headers*).
 - Contains metadata describing program components (*section headers*).

ELF Headers

readelf -a ./cat

```
seed@seed-vm ~/Programs> readelf -a ./cat
ELF Header:
  Magic:   7f 45 4c 46 02 01 01 00 00 00 00 00 00 00 00 00
  Class:                               ELF64
  Data:                                   2's complement, little endian
  Version:                               1 (current)
  OS/ABI:                                UNIX - System V
  ABI Version:                           0
  Type:                                  DYN (Position-Independent Executable file)
  Machine:                               AArch64
  Version:                               0x1
  Entry point address:                   0x7c0
  Start of program headers:               64 (bytes into file)
  Start of section headers:              7240 (bytes into file)
  Flags:                                  0x0
  Size of this header:                    64 (bytes)
  Size of program headers:                56 (bytes)
  Number of program headers:               13
  Size of section headers:                64 (bytes)
  Number of section headers:               28
  Section header string table index:      20

Section Headers:
  [Nr] Name              Type              Address              Offset
       Size              EntSize          Flags  Link  Info  Align
  [ 0]                  NULL              0000000000000000    00000000
```


ELF Program Headers

- Program headers specify *segments* that will be loaded to memory for executing the program. Most important entry types:
 - **INTERP:** defines the library that should be used to load this ELF into memory.
 - **LOAD:** defines a part of the file that should be loaded into memory.
- PH are *the* source of information used when loading a file. It tells the OS how to create a process image
 - Always exists for any valid ELF.

ELF Program Headers

readelf -a ./cat

Program Headers:

Type	Offset FileSiz	VirtAddr MemSiz	PhysAddr Flags Align
PHDR	0x0000000000000040	0x0000000000000040	0x0000000000000040
	0x00000000000002d8	0x00000000000002d8	R 0x8
LOAD	0x0000000000000000	0x0000000000000000	0x0000000000000000
	0x0000000000000aa8	0x0000000000000aa8	R E 0x10000
GNU_STACK	0x0000000000000000	0x0000000000000000	0x0000000000000000
	0x0000000000000000	0x0000000000000000	RW 0x10
GNU_EH_FRAME	0x00000000000009bc	0x00000000000009bc	0x00000000000009bc
	0x000000000000003c	0x000000000000003c	R 0x4
LOAD	0x0000000000000d60	0x0000000000010d60	0x0000000000010d60
	0x00000000000002b0	0x00000000000002b8	RW 0x10000
GNU_RELRO	0x0000000000000d60	0x0000000000010d60	0x0000000000010d60
	0x00000000000002a0	0x00000000000002a0	R 0x1
LOAD	0x0000000000010000	0x0000000000020000	0x0000000000020000
	0x00000000000001c8	0x00000000000001c8	RW 0x10000
NOTE	0x0000000000010180	0x0000000000020180	0x0000000000020180
	0x0000000000000020	0x0000000000000020	R 0x4
NOTE	0x00000000000101a0	0x00000000000201a0	0x00000000000201a0
	0x0000000000000024	0x0000000000000024	R 0x4
INTERP	0x0000000000020000	0x0000000000030000	0x0000000000030000
	0x000000000000001b	0x000000000000001b	R 0x1
[Requesting program interpreter: /lib/ld-linux-aarch64.so.1]			
LOAD	0x0000000000020000	0x0000000000030000	0x0000000000030000
	0x0000000000000020	0x0000000000000020	RW 0x10000
DYNAMIC	0x0000000000030000	0x0000000000040000	0x0000000000040000
	0x0000000000000210	0x0000000000000210	RW 0x8
LOAD	0x0000000000030000	0x0000000000040000	0x0000000000040000
	0x00000000000002f8	0x00000000000002f8	RW 0x10000

0x1b=27

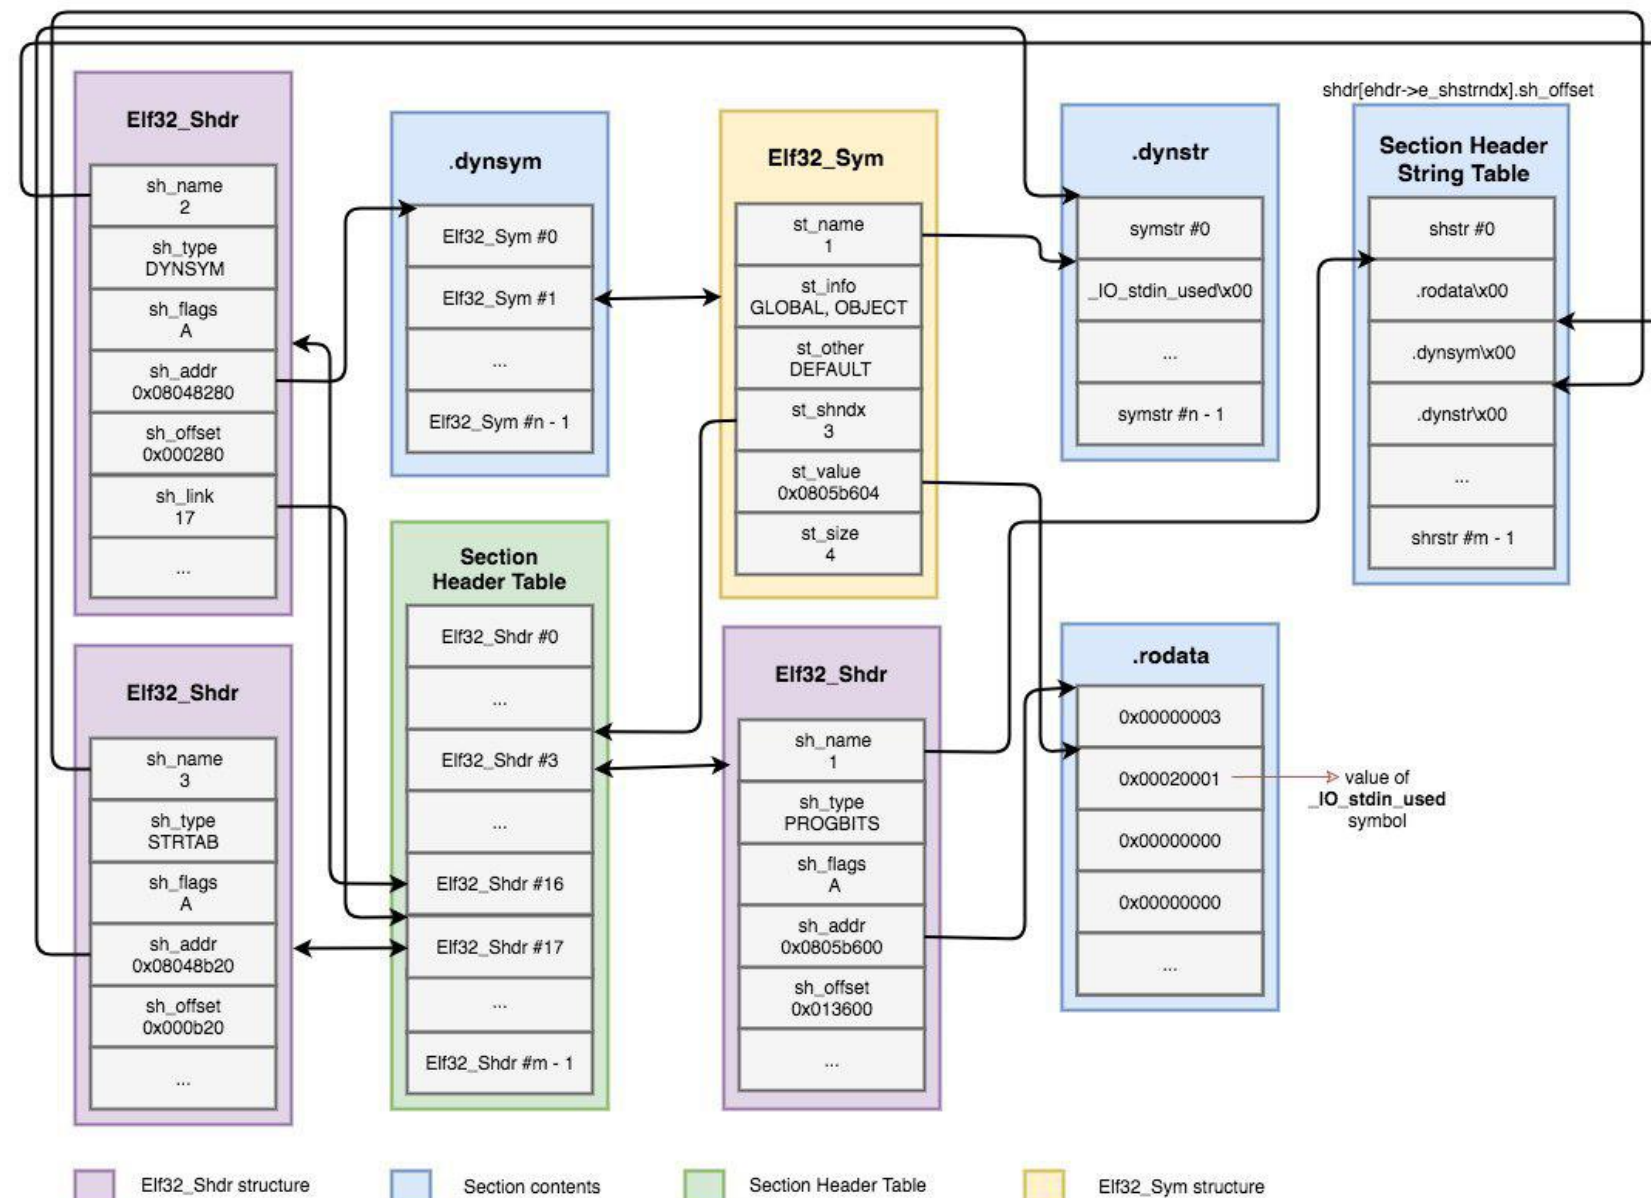
27 bytes
(including the
end '\0')

ELF Section Headers

- A different view of the ELF with useful information for introspection, debugging, etc.
- Important sections:
 - **.text**: the executable code of your program.
 - **.plt** and **.got**: used to resolve and dispatch library calls.
 - **.data**: used for *pre-initialized global writable* data (such as global arrays with initial values)
 - **.rodata**: used for *global read-only* data (such as string constants)
 - **.bss**: used for uninitialized *global writable* data (such as global arrays without initial values)
- Section headers are not a necessary part of the ELF: only *segments* (defined via program headers) are needed for loading and operation! Section headers are just metadata.

Symbols

Binaries (and libraries) that use **dynamically** loaded libraries rely on symbols (names) to find libraries, resolve function calls into those libraries, etc.




Interacting with ELF

- Several ways to dig in:
 - **gcc** to make your ELF.
 - **readelf** to parse the ELF header.
 - **objdump** to parse the ELF header and disassemble the source code.
 - **nm** to view your ELF's symbols.
 - **patchelf** to change some ELF properties.
 - **objcopy** to swap out ELF sections.
 - **strip** to remove otherwise-helpful information (such as symbols).
 - **kaitai** struct (<https://ide.kaitai.io/>) to look through your ELF interactively.

Linux Process Loading and Execution

From a C program to a process



The life of a cat

```
$ cat cat.c
// implements a cat!
int main(int argc, char **argv) {
    char buf[2048];
    int n;
    int fd = argc == 1 ? 0: open(argv[1], 0);
    while ((n = read(fd, buf, 2048)) > 0 && write(1, buf, n) > 0);
}
```

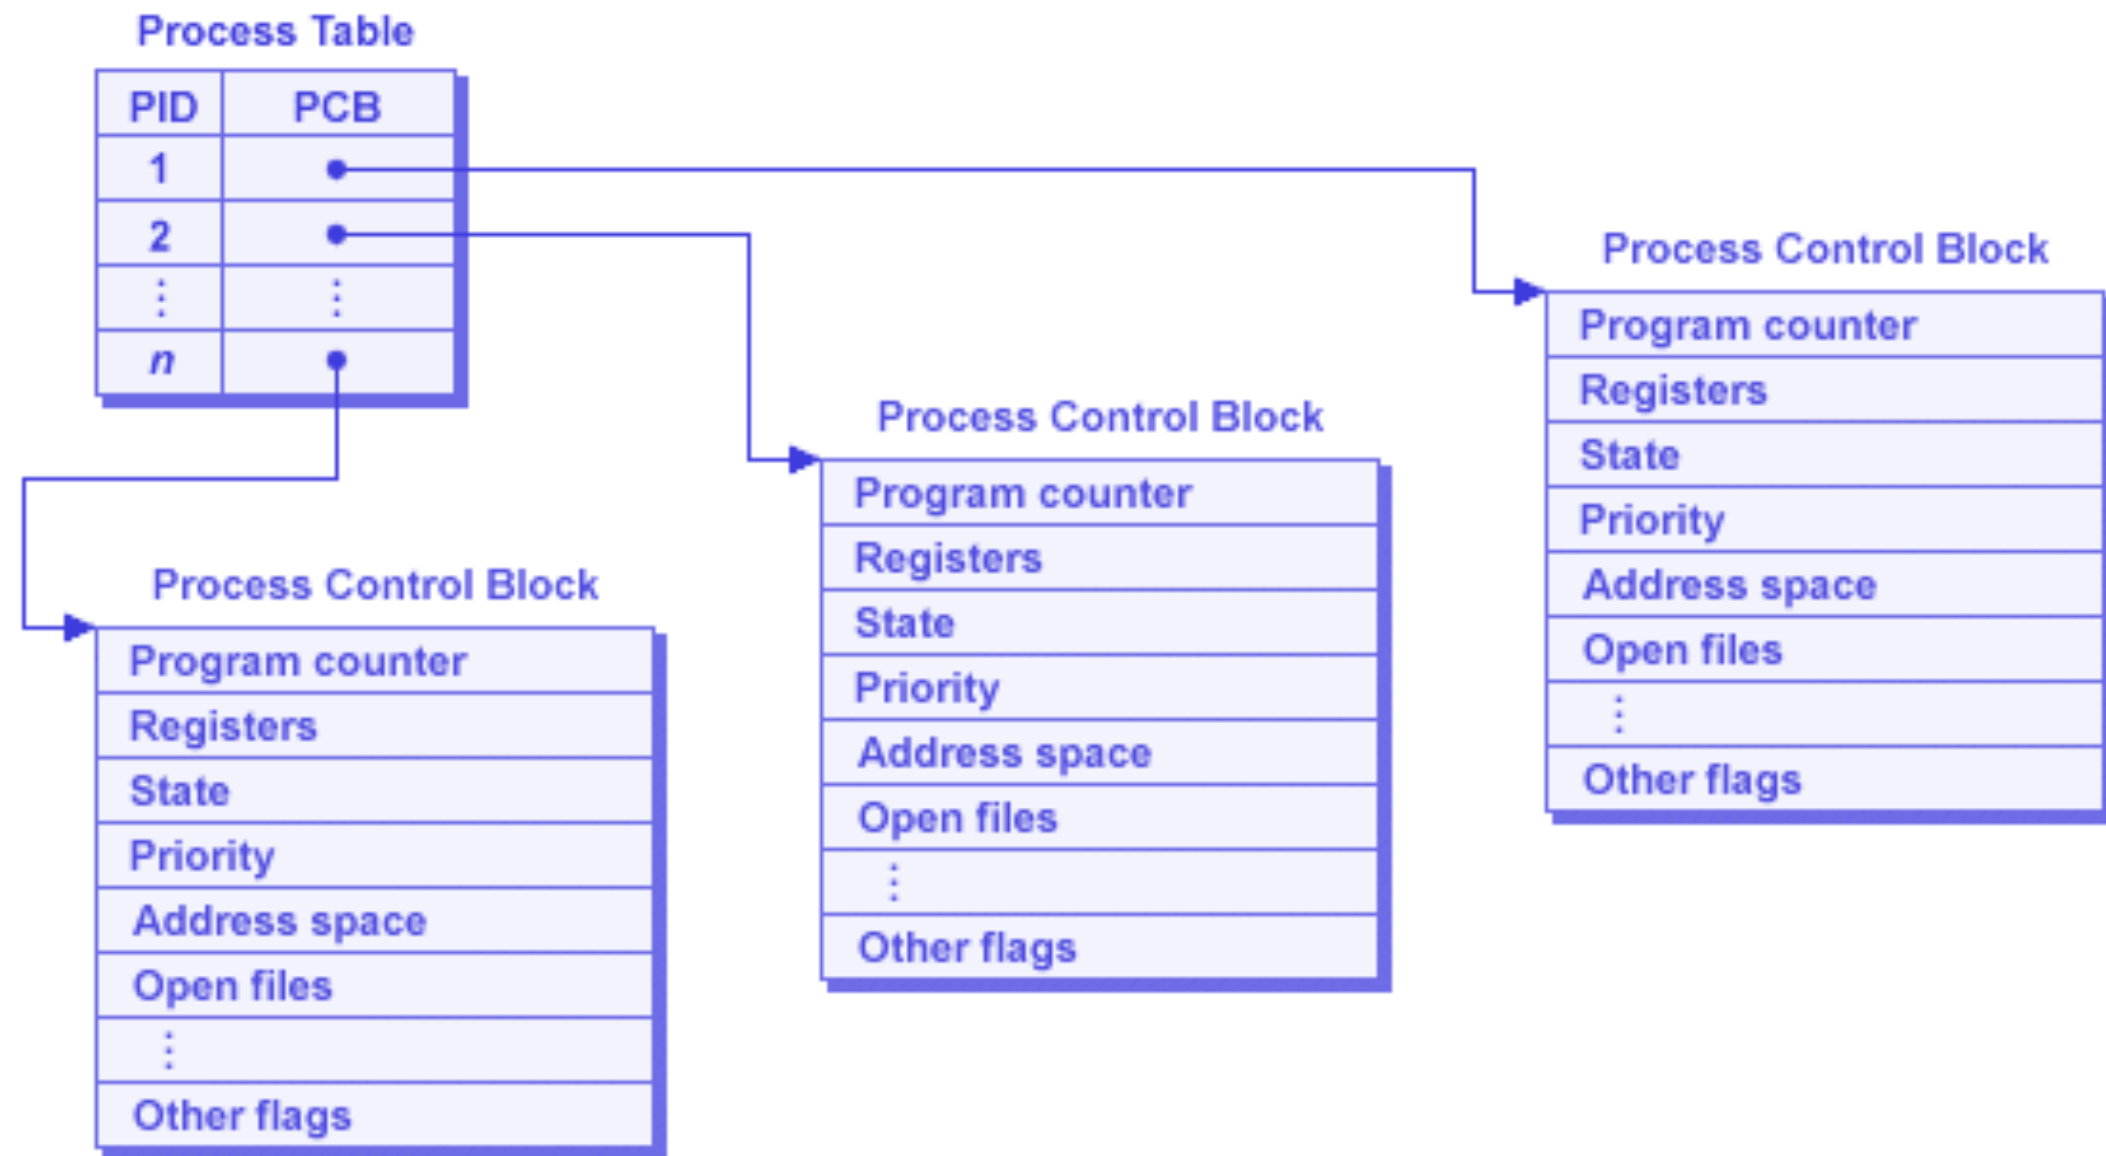
1. A process is created
2. cat is loaded
3. cat is initialized
4. cat launched
5. cat reads args and envs
6. cat does its thing
7. cat terminates

A process is created

- Every Linux process has
 - Process ID (PID)
 - Current state (running, ready, waiting, etc.)
 - Program counter (indicating the next instruction to execute)
 - Memory allocation details: virtual memory space
 - I/O status and other resources (files, pipes, sockets)
 - **Security context:** effective UID/GID, saved UID/GID, capabilities


A process is created

Process Control Block (PCB)



A process is created

- In Linux, processes propagate by mitosis
 - **fork** and (more recently) **clone** are [system calls](#) that create a nearly exact copy of the calling process: a parent and a child
 - The child then uses **execve** [syscall](#) to replace itself with another process
- Example
 - Type **./cat** in bash
 - bash **forks** itself into the old parent and child process
 - the child **execves** **./cat**, becoming **./cat**



(Before) Loading Cat

- Before anything is loaded, the kernel does some validations
 - permission, memory requirements, etc
 - If failed, then **execve** fails
- Then the OS sets up a new process for the program to run in, including a virtual address space.
- Then the OS tries to find the [interpreter](#) for the program

Loading Cat

- If file starts with `#!`,
 - kernel extracts the interpreter from the rest of that line and executes this interpreter **with the original file as an argument**
- E.g. shell script
- If file matches a format (**magic number**) in `/proc/sys/fs/binfmt_misc`
 - kernel executes the interpreter specified for that format **with the original file as an argument**

```
seed@seed-vm ~/Programs> cat /proc/sys/fs/binfmt_misc/python3.10
enabled
interpreter /usr/bin/python3.10
flags:
offset 0
magic 6f0d0d0a
```

Loading Cat

- If file is a **dynamically-linked ELF** (Executable and Linkable Format)
 - ▶ kernel reads the interpreter/loader from the ELF,
 - ▶ kernel loads the interpreter and the original file
 - ▶ The interpreter takes control
- If file is a **statically-linked ELF**
 - kernel loads it
- This can be recursive

Loading a Dynamically-linked Cat

The interpreter

Process loading is done by the ELF interpreter specified in the binary.

```
seed@seed-vm ~/Programs> readelf -a ./cat | grep interpreter  
[Requesting program interpreter: /lib/ld-linux-aarch64.so.1]
```

Colloquially known as "the loader".

Can be overridden:

```
seed@seed-vm ~/Programs> /lib/ld-linux-aarch64.so.1 ./cat cat.c  
int main(int argc, char **argv) {  
    char buf[2048];  
    int n;  
    int fd = argc == 1 ? 0: open(argv[1], 0);  
  
    while ((n = read(fd, buf, 2048)) > 0 && write(1, buf, n) > 0);  
}
```

Or changed permanently:

```
seed@seed-vm ~/Programs> patchelf --set-interpreter /some/interp ./cat  
seed@seed-vm ~/Programs> readelf -a ./cat | grep interpreter  
[Requesting program interpreter: /some/interp]
```

Loading a Dynamically-linked Cat

The Loading Process

- The program and its interpreter are loaded by the kernel.
- The interpreter locates the libraries.
 - `LD_PRELOAD` environment variable, and anything in `/etc/ld.so.preload`
 - `LD_LIBRARY_PATH` environment variable (can be set in the shell)
 - `DT_RUNPATH` or `DT_RPATH` specified in the binary file (both can be modified with `patchelf`)
 - system-wide configuration (`/etc/ld.so.conf`)
 - `/lib` and `/usr/lib`
- The interpreter loads the libraries.
 - These libraries can depend on other libraries, causing more to be loaded
 - Relocations updated

Loading a Dynamically-linked Cat

The Loading Process

```
seed@seed-vm ~/Programs> strace ./cat cat.c
execve("./cat", ["/cat", "cat.c"], 0xfffffa344c68 /* 43 vars */) = 0
brk(NULL) = 0xafad3790e000
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xf7d3e7ac8000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=58719, ...}, AT_EMPTY_PATH) = 0
mmap(NULL, 58719, PROT_READ, MAP_PRIVATE, 3, 0) = 0xf7d3e7a84000
close(3) = 0
openat(AT_FDCWD, "/lib/aarch64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC) = 3
read(3, "\177ELF\2\1\1\3\0\0\0\0\0\0\0\0\0\0\3\0\267\0\1\0\0\0\340u\2\0\0\0\0"... , 832) = 832
newfstatat(3, "", {st_mode=S_IFREG|0755, st_size=1637400, ...}, AT_EMPTY_PATH) = 0
mmap(NULL, 1805928, PROT_NONE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xf7d3e78cb000
mmap(0xf7d3e78d0000, 1740392, PROT_READ|PROT_EXEC, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0) = 0xf7d3e78d0000
munmap(0xf7d3e78cb000, 20480) = 0
munmap(0xf7d3e7a79000, 44648) = 0
mprotect(0xf7d3e7a58000, 61440, PROT_NONE) = 0
mmap(0xf7d3e7a67000, 24576, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_DENYWRITE, 3, 0x187000) = 0xf7d3e7a67000
mmap(0xf7d3e7a6d000, 48744, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_FIXED|MAP_ANONYMOUS, -1, 0) = 0xf7d3e7a6d000
close(3) = 0
```

Loading a Dynamically-linked Cat

The Loading Process

Change the `LD_LIBRARY_PATH`

```
seed@seed-vm ~/Programs> strace -E LD_LIBRARY_PATH=/some/library/path ./cat cat.c
execve("./cat", ["./cat", "cat.c"], 0xbcd0ad4f46c0 /* 44 vars */) = 0
brk(NULL)                               = 0xbd6747bd2000
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xfd3b6852d000
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/tls/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/tls/aarch64/atomics", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/tls/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/tls/aarch64", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/tls/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/tls/atomics", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/tls/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/tls", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/aarch64/atomics", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/aarch64", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path/atomics", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/some/library/path/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
newfstatat(AT_FDCWD, "/some/library/path", 0xffffffffb52a5b0, 0) = -1 ENOENT (No such file or directory)
openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
newfstatat(3, "", {st_mode=S_IFREG|0644, st_size=58719, ...}, AT_EMPTY_PATH) = 0
```


Loading a Dynamically-linked Cat

The Loading Process

Change both the `LD_LIBRARY_PATH` and `DT_RPATH`

```
seed@seed-vm ~/Programs> patchelf --set-rpath /some/rpath ./cat
seed@seed-vm ~/Programs> strace -E LD_LIBRARY_PATH=/some/library/path ./cat cat.c
```

```
execve("./cat", [".cat", "cat.c"], 0xc21d1982a6c0 /* 44 vars */) = 0
```

```
brk(NULL) = 0xc244ec889000
```

```
mmap(NULL, 8192, PROT_READ|PROT_WRITE, MAP_PRIVATE|MAP_ANONYMOUS, -1, 0) = 0xe9b8cb92f000
```

```
faccessat(AT_FDCWD, "/etc/ld.so.preload", R_OK) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/tls/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/tls/aarch64/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/tls/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/tls/aarch64", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/tls/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/tls/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/tls/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/tls", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/aarch64/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/aarch64", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/library/path/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/library/path", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/tls/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/tls/aarch64/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/tls/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/tls/aarch64", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/tls/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/tls/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/tls/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/tls", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/aarch64/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/aarch64/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/aarch64/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/aarch64", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/atomics/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```

```
newfstatat(AT_FDCWD, "/some/rpath/atomics", 0xffffeef87330, 0) = -1 ENOENT (No such file or directory)
```

```
openat(AT_FDCWD, "/some/rpath/libc.so.6", O_RDONLY|O_CLOEXEC) = -1 ENOENT (No such file or directory)
```


search `LD_LIBRARY_PATH`

search `RPATH`

Loading a Dynamically-linked Cat

Where is cat loaded to?

- Each Linux process has a virtual memory space. It contains:
 - The binary
 - The libraries
 - The "heap" (for dynamically allocated memory)
 - The "stack" (for function local variables)
 - Any memory specifically mapped by the program
 - Some helper regions
 - Kernel code in the "upper half" of memory (above `0x8000000000000000` on 64-bit architectures), inaccessible to the process
- Virtual memory is dedicated to your process.
Physical memory is shared among the whole system.
- You can see this whole space by looking at `/proc/self/maps`



Loading a Dynamically-linked Cat

Where is cat loaded to?

```
./cat /proc/self/maps
```

```
seed@seed-vm ~/Programs> ./cat /proc/self/maps
```

b0b6c4950000-b0b6c4951000	r-xp	00000000	103:02	428496	/home/seed/Programs/cat	cat's binary
b0b6c4960000-b0b6c4961000	r--p	00000000	103:02	428496	/home/seed/Programs/cat	
b0b6c4961000-b0b6c4962000	rw-p	00001000	103:02	428496	/home/seed/Programs/cat	
b0b6c4970000-b0b6c4971000	rw-p	00010000	103:02	428496	/home/seed/Programs/cat	
b0b6c4980000-b0b6c4981000	rw-p	00020000	103:02	428496	/home/seed/Programs/cat	
b0b6c4990000-b0b6c4991000	rw-p	00030000	103:02	428496	/home/seed/Programs/cat	
e5881e1f0000-e5881e378000	r-xp	00000000	103:02	1053664	/usr/lib/aarch64-linux-gnu/libc.so.6	libc
e5881e378000-e5881e387000	---p	00188000	103:02	1053664	/usr/lib/aarch64-linux-gnu/libc.so.6	
e5881e387000-e5881e38b000	r--p	00187000	103:02	1053664	/usr/lib/aarch64-linux-gnu/libc.so.6	
e5881e38b000-e5881e38d000	rw-p	0018b000	103:02	1053664	/usr/lib/aarch64-linux-gnu/libc.so.6	
e5881e38d000-e5881e399000	rw-p	00000000	00:00	0		
e5881e3ac000-e5881e3d7000	r-xp	00000000	103:02	1053335	/usr/lib/aarch64-linux-gnu/ld-linux-aarch64.so.1	linker/loader
e5881e3df000-e5881e3e3000	rw-p	00000000	00:00	0		
e5881e3e3000-e5881e3e5000	r--p	00000000	00:00	0	[vvar]	
e5881e3e5000-e5881e3e6000	r-xp	00000000	00:00	0	[vdso]	
e5881e3e6000-e5881e3e8000	r--p	0002a000	103:02	1053335	/usr/lib/aarch64-linux-gnu/ld-linux-aarch64.so.1	
e5881e3e8000-e5881e3ea000	rw-p	0002c000	103:02	1053335	/usr/lib/aarch64-linux-gnu/ld-linux-aarch64.so.1	
ffffff630d000-ffffff632e000	rw-p	00000000	00:00	0	[stack]	

The Standard C Library

- `libc.so` is linked by almost every process.
- Provides functionality you take for granted:
 - `printf()`
 - `scanf()`
 - `socket()`
 - `atoi()`
 - `malloc()`
 - `free()`
 - ...

Loading a Statically-linked Cat


The binary is loaded by the kernel. DONE.

We can force linking a library statically:

► `gcc -static -o cat-stat cat.c`

...and we end up with a much bigger cat:

```
seed@seed-vm ~/Programs> ll cat cat-stat  
-rwxrwxr-x 1 seed seed 8.9K Nov  7 23:09 cat  
-rwxrwxr-x 1 seed seed 632K Nov  7 23:09 cat-stat
```



Loading a Statically-linked Cat

The binary is loaded by the kernel. DONE.

We can force linking a library statically:

► `gcc -static -o cat-stat cat.c`

...and we end up with a much bigger cat, with no `libc.so` linked

```
seed@seed-vm ~/Programs> ./cat-stat /proc/self/maps
00400000-0047e000 r-xp 00000000 103:02 404267 /home/seed/Programs/cat-stat
0048d000-00491000 r--p 0007d000 103:02 404267 /home/seed/Programs/cat-stat
00491000-00494000 rw-p 00081000 103:02 404267 /home/seed/Programs/cat-stat
00494000-00499000 rw-p 00000000 00:00 0
282bb000-282dd000 rw-p 00000000 00:00 0 [heap]
fbc6b9e8e000-fbc6b9e90000 r--p 00000000 00:00 0 [vvar]
fbc6b9e90000-fbc6b9e91000 r-xp 00000000 00:00 0 [vdso]
ffffc166a000-ffffc168b000 rw-p 00000000 00:00 0 [stack]
```


Launch Cat

- A normal ELF automatically calls `__libc_start_main()` in `libc`, which in turn calls the program's `main()` function.

```
seed@seed-vm ~/Programs> readelf -a ./cat | grep Entry
Entry point address: 0x7c0
```

```
seed@seed-vm ~/Programs> objdump -d -j .text ./cat | head -n 40
```

```
./cat:      file format elf64-littleaarch64
```

```
Disassembly of section .text:
```


```
000000000000007c0 <.text>:
7c0: d503201f      nop
7c4: d280001d      mov     x29, #0x0           // #0
7c8: d280001e      mov     x30, #0x0           // #0
7cc: aa0003e5      mov     x5, x0
7d0: f94003e1      ldr     x1, [sp]
7d4: 910023e2      add     x2, sp, #0x8
7d8: 910003e6      mov     x6, sp
7dc: 90000080      adrp    x0, 10000 <read@plt+0xf870>
7e0: f947f800      ldr     x0, [x0, #4080]
7e4: d2800003      mov     x3, #0x0           // #0
7e8: d2800004      mov     x4, #0x0           // #0
7ec: 97ffffcd      bl      720 <__libc_start_main@plt>
7f0: 97ffffe4      bl      780 <abort@plt>
```

- Your code is running! Now what?

Cat reads its args and envs

- `int main(int argc, void **argv, void **envp);`
- Your process's entire input from the outside world, at launch, comprises of:
 - The loaded objects (binaries and libraries)
 - command-line arguments in **argv**
 - "environment" in **envp**
- Of course, processes need to keep interacting with the outside world.

Cat does its thing



Using library functions

- The binary's import symbols have to be resolved using the libraries' export symbols.
- In the past, this was an *on-demand* process and carried great peril.
- In modern times, this is all done when the binary is *loaded*, and is much safer.
- We'll explore this further in the future.

Using library functions

```
seed@seed-vm ~/Programs> nm ./cat
0000000000000278 r __abi_tag
                U abort@GLIBC_2.17
00000000000011018 B __bss_end__
00000000000011018 B __bss_end__
00000000000011010 B __bss_start
00000000000011010 B __bss_start__
00000000000007f4 t call_weak_fn
00000000000011010 b completed.0
                w __cxa_finalize@GLIBC_2.17
00000000000011000 D __data_start
00000000000011000 W data_start
00000000000000810 t deregister_tm_clones
00000000000000880 t __do_global_dtors_aux
00000000000010d68 d __do_global_dtors_aux_fini_array_entry
00000000000011008 D __dso_handle
00000000000010d70 a _DYNAMIC
00000000000011010 D _edata
00000000000011018 B __end__
00000000000011018 B _end
000000000000009a4 T _fini
000000000000008d0 t frame_dummy
00000000000010d60 d __frame_dummy_init_array_entry
00000000000000aa4 r __FRAME_END__
00000000000010fc8 a _GLOBAL_OFFSET_TABLE_
                w __gmon_start__
000000000000009bc r __GNU_EH_FRAME_HDR
000000000000006e8 T _init
000000000000009b8 R _IO_stdin_used
                w _ITM_deregisterTMCloneTable
                w _ITM_registerTMCloneTable
                U __libc_start_main@GLIBC_2.34
000000000000008d4 T main
                U open@GLIBC_2.17
                U read@GLIBC_2.17
00000000000000840 t register_tm_clones
                U __stack_chk_fail@GLIBC_2.17
                U __stack_chk_guard@GLIBC_2.17
000000000000007c0 t _start
00000000000011010 D __TMC_END__
                U write@GLIBC_2.17
```

U: undefined symbols

will be validated by the linker and loaded by the loader

Interacting with the environment

- Almost all programs have to interact with the outside world!
- This is primarily done via **system calls** (`man syscalls`). Each system call is well-documented in section 2 of the man pages (i.e., `man 2 open`).
- <https://x86.syscall.sh/>
- We can trace process system calls using **strace**.

System Calls

- System calls have very well-defined interfaces that very rarely change.
- There are over 300 system calls in Linux. Here are some examples:
 - `int open(const char *pathname, int flags)` - returns a file new file descriptor of the open file (also shows up in `/proc/self/fd!`)
 - `ssize_t read(int fd, void *buf, size_t count)` - reads data from the file descriptor
 - `ssize_t write(int fd, void *buf, size_t count)` - writes data to the file descriptor
 - `pid_t fork()` - forks off an identical child process. Returns 0 if you're the child and the PID of the child if you're the parent.
 - `int execve(const char *filename, char **argv, char **envp)` - replaces your process.
 - `pid_t wait(int *wstatus)` - wait child termination, return its PID, write its status into `*wstatus`.
 - `long syscall(long syscall, ...)` - invoke specified `syscall`.
- Typical signal combinations:
 - `fork, execve, wait` (think: a shell)
 - `open, read, write` (cat)

Signals

- System calls are a way for a process to call into the OS. What about the other way around?
- Enter: signals. Relevant system calls:
 - `sighandler_t signal(int signum, sighandler_t handler)` - register a **signal handler**
 - `int sigaction(int signum, const struct sigaction *act, struct sigaction *oldact)` - more modern way of registering a signal handler
 - `int kill(pid_t pid, int sig)` - **send** a signal to a process.
- Signals **pause** process execution and **invoke** the handler.
- Handlers are functions that take one argument: the signal number.
- Without a handler for a signal, the default action is used (often, kill).
- **SIGKILL** (signal 9) and **SIGSTOP** (signal 19) cannot be handled.

Signals

Full list in section 7 of signal manual: `man 7 signal`

Signal	Standard	Action	Comment
SIGABRT	P1990	Core	Abort signal from abort(3)
SIGALRM	P1990	Term	Timer signal from alarm(2)
SIGBUS	P2001	Core	Bus error (bad memory access)
SIGCHLD	P1990	Ign	Child stopped or terminated
SIGCLD	-	Ign	A synonym for SIGCHLD
SIGCONT	P1990	Cont	Continue if stopped
SIGEMT	-	Term	Emulator trap
SIGFPE	P1990	Core	Floating-point exception
SIGHUP	P1990	Term	Hangup detected on controlling terminal or death of controlling process
SIGILL	P1990	Core	Illegal Instruction
SIGINFO	-		A synonym for SIGPWR
SIGINT	P1990	Term	Interrupt from keyboard
SIGIO	-	Term	I/O now possible (4.2BSD)
SIGIOT	-	Core	IOT trap. A synonym for SIGABRT
SIGKILL	P1990	Term	Kill signal
SIGLOST	-	Term	File lock lost (unused)
SIGPIPE	P1990	Term	Broken pipe: write to pipe with no readers; see pipe(7)
SIGPOLL	P2001	Term	Pollable event (Sys V); synonym for SIGIO
SIGPROF	P2001	Term	Profiling timer expired
SIGPWR	-	Term	Power failure (System V)

Shared memory

- Another way of interacting with the outside world is by sharing memory with other processes.
- Requires system calls to establish, but once established, communication happens without system calls.
 - `int shm_fd = shm_open("/myshm", O_CREAT | O_RDWR, 0666);`
 - `char *ptr = (char*) mmap(0, 1024, PROT_WRITE, MAP_SHARED, shm_fd, 0);` // map shared memory to process' own mem address
 - `strncpy(ptr, 1024, "hello");` // write data to shared memory
- Easy way: use a shared memory-mapped file in `/dev/shm`.
 - ▶ `FILE *fp = fopen("/dev/shm/my_shm", "r+");`
 - ▶ `fprintf(fp, "Writing to the shared mem\n");`
 - ▶ `fgets(buf, sizeof(buf), fp)`

Cat terminates

- Processes terminate by one of two ways:
 1. Receiving an unhandled signal.
 2. Calling the `exit()` system call: `int exit(int status)`
- All processes must be "reaped":
 - After termination, they will remain in a zombie state until they are `wait()`ed on by their parent.
 - When this happens, their exit code will be returned to the parent, and the process will be freed.
 - If their parent dies without `wait()`ing on them, they are re-parented to PID 1 and will stay there until they're cleaned up.

Questions?